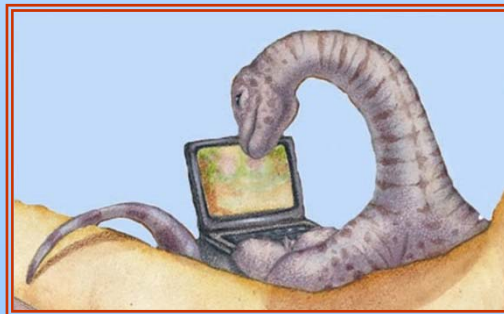


Chapter 5: CPU Scheduling





Chapter 5: CPU Scheduling

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Multiple-Processor Scheduling
- Real-Time Scheduling
- Thread Scheduling
- Operating Systems Examples





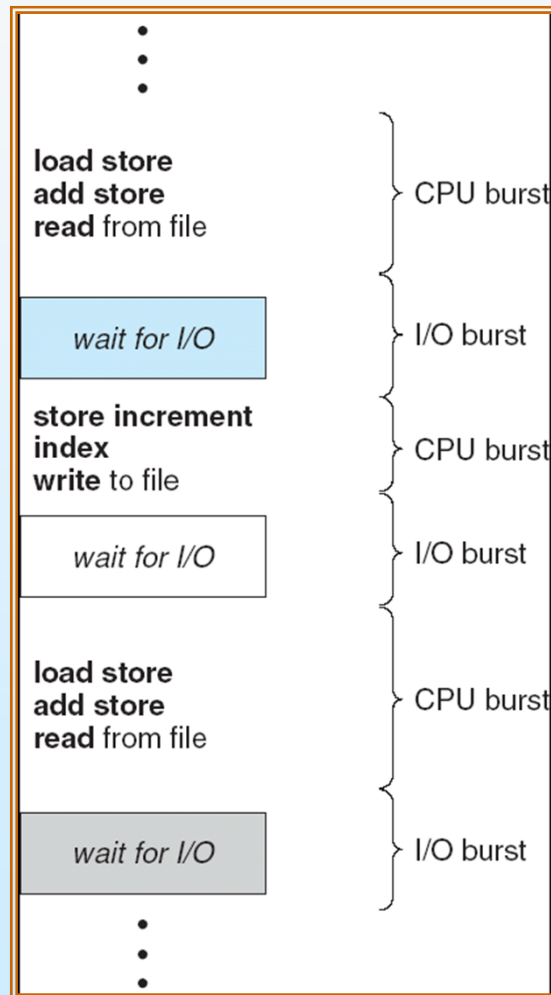
Basic Concepts

- **Maximum CPU** utilization obtained with multiprogramming
- **CPU-I/O** Burst Cycle – Process execution consists of a ***cycle*** of CPU execution and I/O wait
- CPU burst distribution






Alternating Sequence of CPU And I/O Bursts





Histogram of CPU-burst Times





CPU Scheduler

- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them
- CPU scheduling decisions may take place when a process:
 1. Switches from **running** to **waiting** state
 2. Switches from **running** to **ready** state
 3. Switches from **waiting** to **ready**
 4. **Terminates**
- Scheduling under 1 and 4 is **nonpreemptive**
- All other scheduling is **preemptive**





Dispatcher

- Dispatcher module **gives control of** the CPU to the process selected **by the short-term scheduler**; this involves:
 - switching context
 - **switching to user mode**
 - **jumping to the proper location in the user program to restart that program**
- ***Dispatch latency*** – time it takes for the dispatcher to **stop one** process and **start another** running





Scheduling Criteria

- **CPU utilization** – keep the CPU **as busy as possible**
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the **first response** is produced, **not** output (for time-sharing environment)





Optimization Criteria

- **Max** CPU utilization
- **Max** throughput
- **Min** turnaround time
- **Min** waiting time
- **Min** response time





First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1 , P_2 , P_3
The Gantt Chart for the schedule is:



- **Waiting time** for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- **Average waiting time:** $(0 + 24 + 27)/3 = 17$





FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- **Waiting time** for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- **Average waiting time:** $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- **Convoy effect** short process behind long process





Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of **its next CPU burst**. Use these lengths to schedule **the process with the shortest time**
- **Two schemes**
 - **Nonpreemptive** – once CPU given to the process it cannot be **preempted** until completes its CPU burst
 - **Preemptive** – if a new process arrives with CPU burst length less than **remaining time of current executing process**, preempt. This scheme is known as the **Shortest-Remaining-Time-First (SRTF)**
- SJF is **optimal** – **gives minimum average waiting time** for a given set of processes






Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SJF (non-preemptive)



- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$






Example of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SJF (preemptive)



- Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$






Determining Length of Next CPU Burst

- Can only estimate the length
- Can be done by using the length of previous CPU bursts, using exponential averaging
 1. t_n = actual length of n^{th} CPU burst
 2. τ_{n+1} = predicted value for the next CPU burst
 3. $\alpha, 0 \leq \alpha \leq 1$
 4. Define :





Prediction of the Length of the Next CPU Burst





Examples of Exponential Averaging

- $\alpha = 0$

- $\tau_{n+1} = \tau_n$
- Recent history does not count

- $\alpha = 1$

- $\tau_{n+1} = \alpha t_n$
- Only the actual last CPU burst counts

- If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$

- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor





Priority Scheduling

- A priority number (integer) is **associated with each process**
- The CPU is allocated to the process with the highest priority
(**smallest integer** \equiv **highest priority**)
 - Preemptive
 - nonpreemptive
- SJF is a priority scheduling **where priority is the predicted next CPU burst time**
- Problem \equiv **Starvation** – low priority processes may never execute
- Solution \equiv **Aging** – as time progresses increase the priority of the process





Round Robin (RR)

- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Performance
 - q large \Rightarrow FIFO
 - q small $\Rightarrow q$ must be large with respect to context switch, otherwise overhead is too high






Example of RR with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24

- The Gantt chart is:




- Typically, higher average turnaround than SJF, but better *response*






Time Quantum and Context Switch Time





Turnaround Time Varies With The Time Quantum





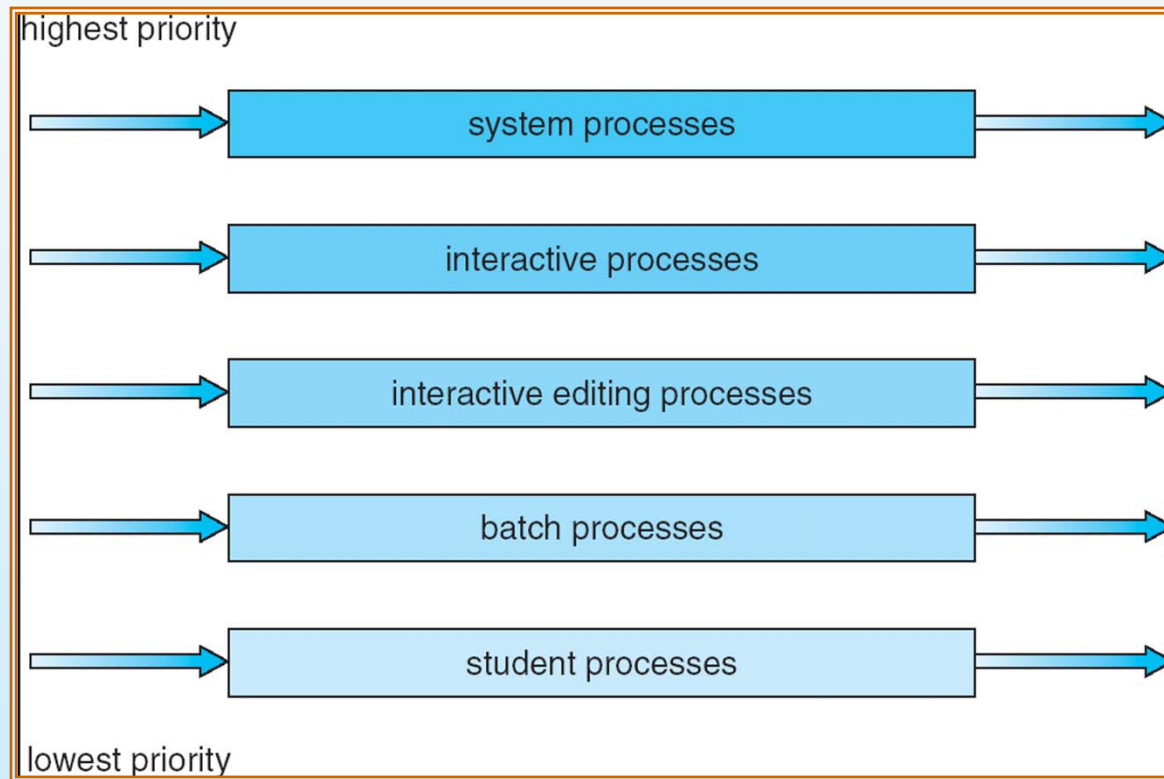
Multilevel Queue

- Ready queue is **partitioned into separate queues**:
 - foreground (interactive)
 - background (batch)
- Each queue has **its own scheduling algorithm**
 - **foreground** – RR
 - **background** – FCFS
- Scheduling must be done between the queues
 - **Fixed priority scheduling**; (i.e., serve all from foreground then from background). **Possibility of starvation**.
 - **Time slice** – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
 - 20% to background in FCFS





Multilevel Queue Scheduling





Multilevel Feedback Queue

- A process can move between the various queues; aging can be implemented this way
- Multilevel-feedback-queue scheduler defined by the following **parameters**:
 - number of queues
 - scheduling algorithms **for each queue**
 - method used to determine when to **upgrade** a process
 - method used to determine when to **demote** a process
 - method used to determine which queue a process will enter when that process needs service






Example of Multilevel Feedback Queue

- Three queues:
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS
- Scheduling
 - A new job enters queue Q_0 which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
 - At Q_1 job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2 .





Multilevel Feedback Queues





Multiple-Processor Scheduling

- CPU scheduling **more complex** when multiple CPUs are available
- *Homogeneous processors* within a multiprocessor
- *Load sharing*
- *Asymmetric multiprocessing* – only one processor accesses the system data structures, alleviating the need for data sharing





Real-Time Scheduling

- **Hard real-time systems** – required to complete a critical task within a guaranteed amount of time
- **Soft real-time computing** – requires that critical processes receive priority over less fortunate ones





Thread Scheduling

- **Local Scheduling** – How the threads library decides which thread to put onto an available LWP
- **Global Scheduling** – How the kernel decides which kernel thread to run next





Windows XP Priorities

	real-time	high	above normal	normal	below normal	idle priority
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1



End of Chapter 5

