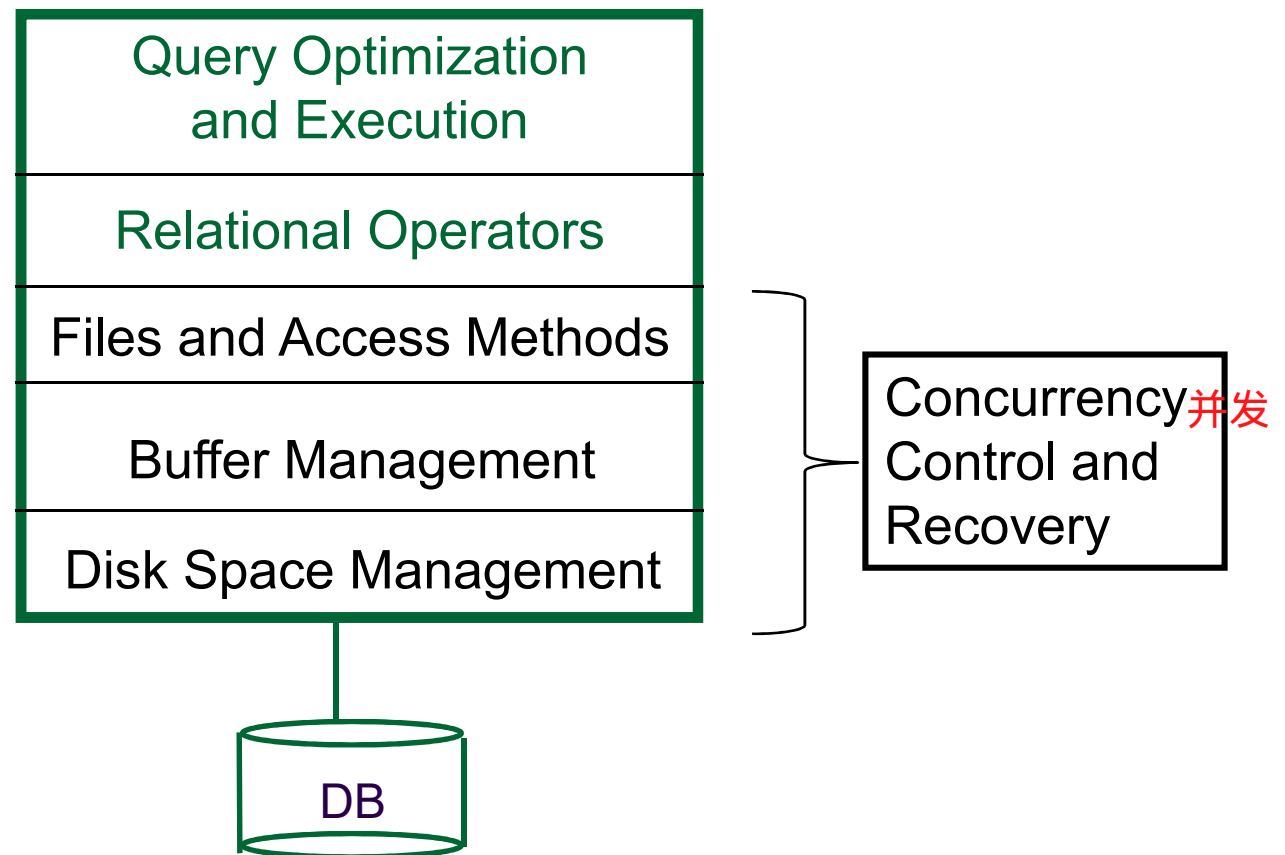


Storing Data: Disks and Files

SUN YAT-SEN UNIVERSITY

Block diagram of a DBMS



Disks

- DBMS stores information on disks.
- Fixed unit of transfer
 - Read/write *disk blocks* or *pages* (8K)
磁盘块 页/数据页/磁盘页
- Not “random access” (vs. RAM)
 - Time to ^{取回}retrieve a block depends on *location*
相对位置
 - *Relative placement* of blocks on disk has major impact on DBMS performance!


Components of a Disk

The platters spin (say, 120 rps).


The arm assembly is moved in or out to position a head on a desired track. Tracks under heads make a *cylinder* (imaginary!).

Only one head reads/writes at any one time.

❖ *Block size* is a multiple of *sector size* (which is fixed).




Accessing a Disk Page



- Time to access (read/write -- 存取时间) a disk block:
 - ❑ *seek time* (寻道时间-moving arms to position disk head on track)
 - ❑ *rotational delay* (旋转延迟-waiting for block to rotate under head)
 - ❑ *transfer time* (传输时间-actually moving data to/from disk surface)
- **Seek time and rotational delay dominate.**
 - ❑ Seek time varies from 0 to 10msec . – 平均时间
 - ❑ Rotational delay varies from 0 to 3msec
 - ❑ Transfer rate around 0.2msec per 8K block
- **Key to lower I/O cost: reduce seek/rotation delays!**
Hardware vs. software solutions?

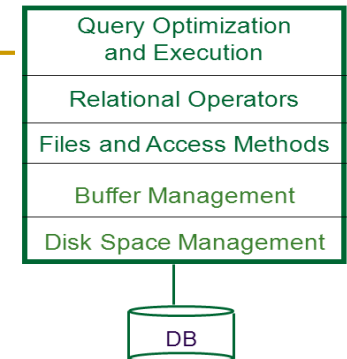
Arranging Pages on Disk



- **`Next'** block concept:
 - ❑ blocks on same track, followed by
 - ❑ blocks on same cylinder, followed by
 - ❑ blocks on adjacent cylinder

- Blocks in a file should be arranged **sequentially** on disk (by `next'), to minimize seek and rotational delay.
 - ❑ For a **sequential scan**, pre-fetching (预读取) several pages at a time is a big win!

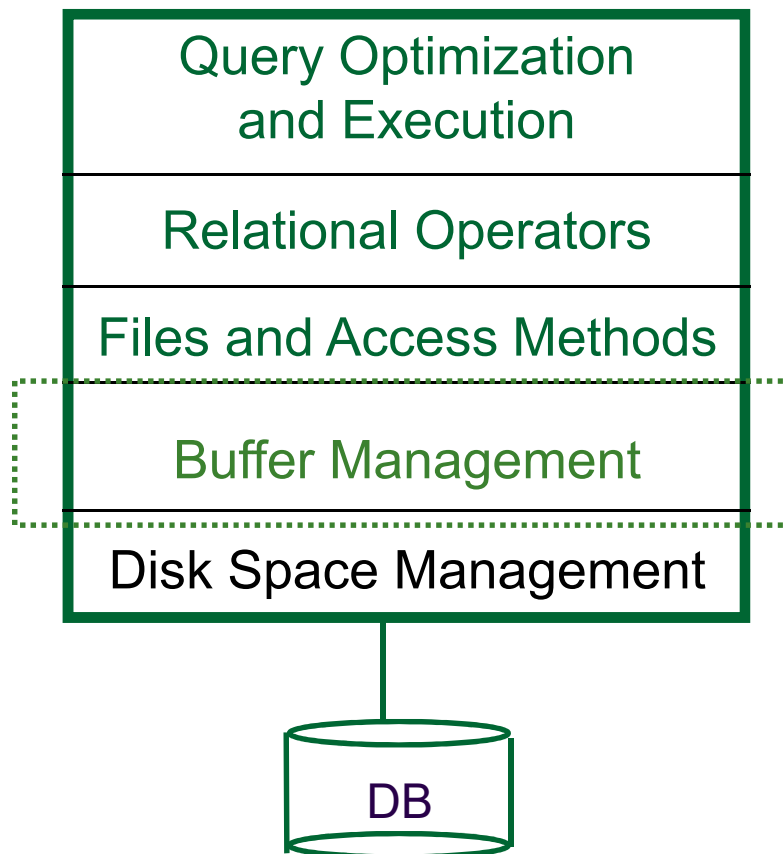
Disk Space Management



- Lowest layer of DBMS, manages space on disk
- 接口: Higher levels call upon this layer to:
 - 分配
 - ❑ allocate/de-allocate a page
 - ❑ read/write a page
- 期望性能: Request for *a sequence of* pages best satisfied by pages *stored sequentially* on disk!
 - ❑ Responsibility of disk space manager.
 - ❑ Higher levels don't know how this is done, or how free space is managed.
 - ❑ They may make performance assumptions!
 - Hence disk space manager should do a decent job.

Context

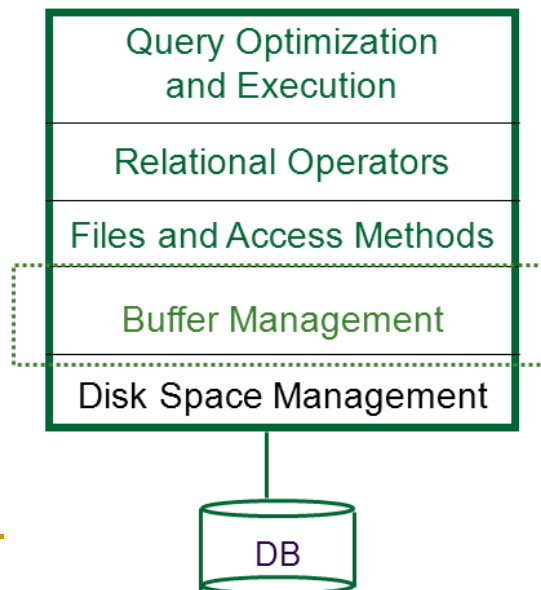
背景



Buffer Management in a DBMS

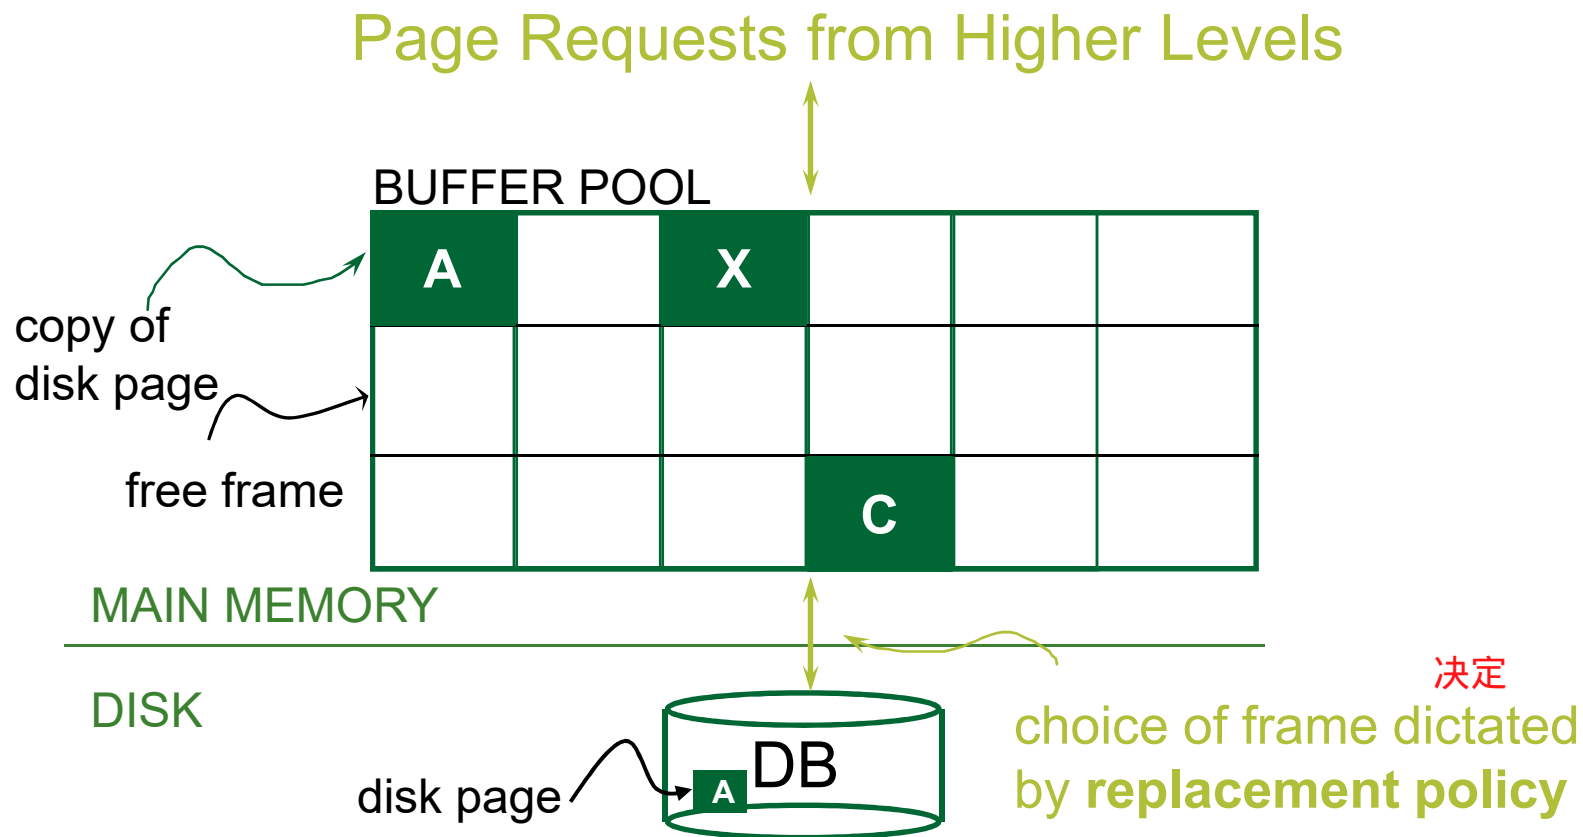
缓冲区管理

- *Data must be in RAM for DBMS to operate on it!*
- *BufMgr hides the fact that not all data is in RAM*




Buffer Management in a DBMS

缓冲区管理—how?




划分帧→(空闲帧/内容帧)→副本→请求→替换策略

When a Page is Requested ...



- Buffer pool information table contains:
<frame#, pageid, pin_count, dirty>
- 1. If requested page is not in pool:
 - a. Choose a frame for *replacement*.
Only “un-pinned” pages are candidates!
 - b. If frame “dirty”, write current page to disk
 - c. Read requested page into frame
- 2. *Pin* the page and return its address.

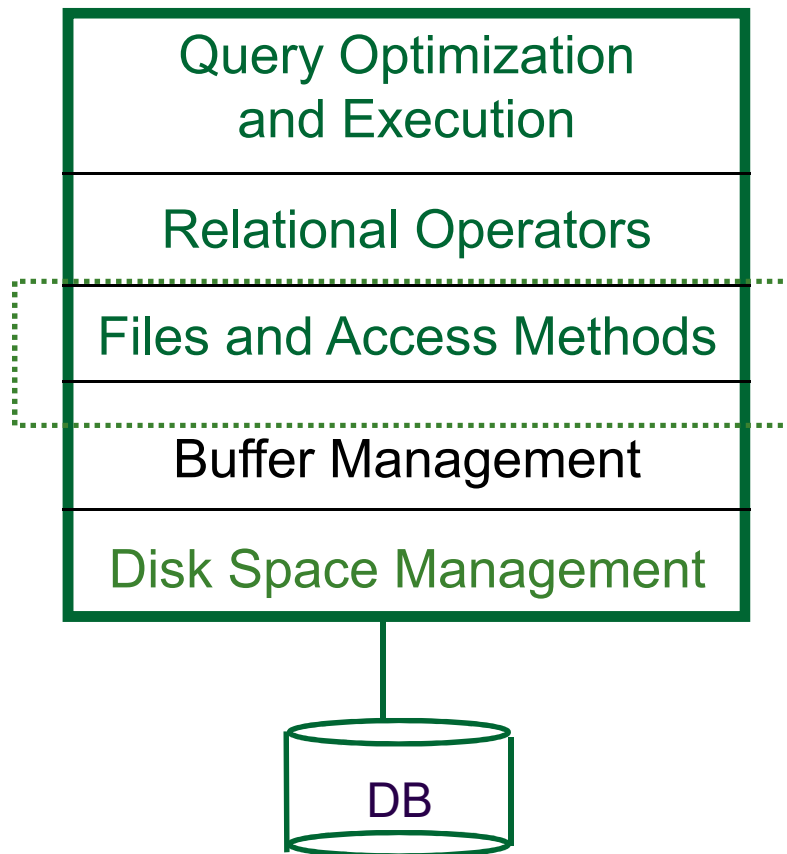
More on Buffer Management



- Requestor of page must eventually:
 1. *unpin* it -- 解钉 `pin_count--`
 2. indicate whether page was modified via *dirty* bit.
表明

<frame#, pageid, pin_count, dirty>
- Page in pool may be requested many times,
 - ❑ a *pin_count* is used.
 - ❑ To pin a page(钉住页): `pin_count++`
 - ❑ A page is a candidate for replacement iff `pin_count == 0` (“unpinned”)

Context




Files of Records (记录文件)

- Blocks are the interface for I/O, but...
 - Higher levels of DBMS operate on *records*, and *files of records*.
 - FILE: A collection of pages, each containing a collection of records. Must support:
 - insert/delete/modify record
 - fetch a particular record (specified using *record id*)
 - scan all records (possibly with some conditions on the records to be retrieved)
-

Unordered (Heap) Files --堆文件


- Collection of records in no particular order.
 - As file ^{收缩}shrinks/grows, disk pages (de)allocated
- To support record level operations, we must:
 - keep track of the *pages* in a file
 - keep track of *free space* on pages
 - keep track of the *records* on a page
- There are many alternatives for keeping track of this.
 - We'll consider *two*.

Heap File Implemented as a List (页链表)



- The header page(首页) id and Heap file name must be stored someplace.
 - Database “catalog”
- Each page contains 2 `pointers' plus data.


Heap File Using a Page Directory(页目录)



- The directory is itself a collection of pages;
 - each page can hold several entries.
 - The entry (目录项) for a page can include the **number** of free bytes on the page. `<pointer,free_number>`
- To **insert a record**, we can search the directory to determine which ~~page has enough space to hold the record.~~

Record Formats(记录格式):

Fixed Length(定长记录)




- Information about field types **same** for all records in a file; stored in *system catalogs*.
- Finding *i'th* field done via arithmetic.

Record Formats:


Variable Length (变长记录)

- Two alternative formats (# fields is fixed):



Fields Delimited by Special Symbols


由特殊符号分隔的字段



Array of Field Offsets-字段偏移量数组

- ✉ Second offers **direct access** to i'th field, efficient storage of nulls (special don't know value); small directory overhead.

Page Formats(页格式): Fixed Length Records



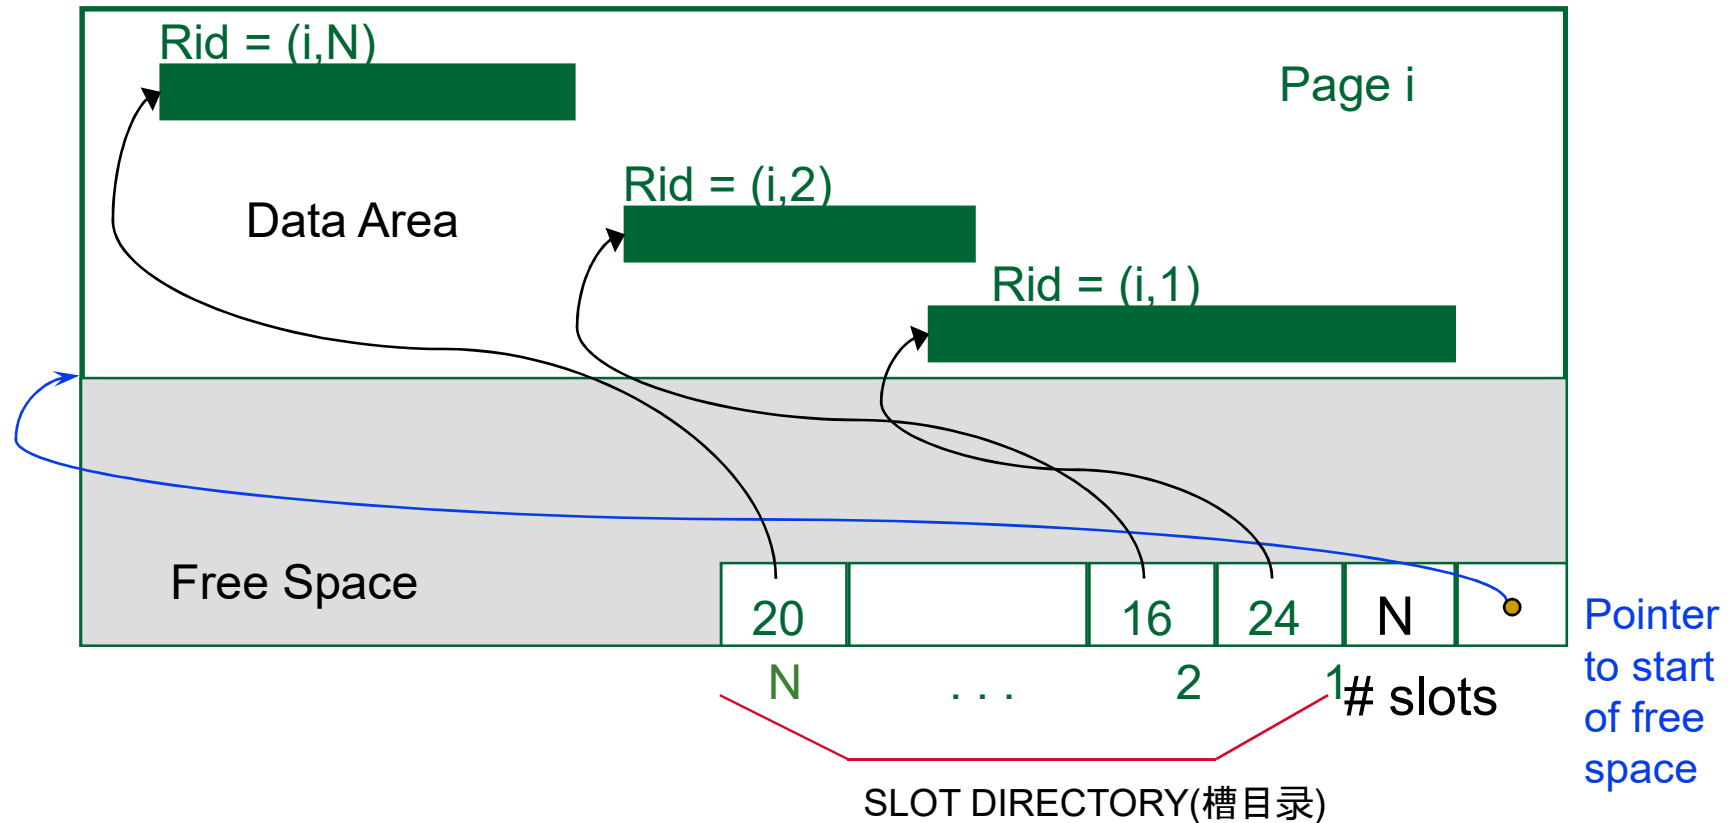
✉ Record id = <page id, slot #>.

✉ In first alternative,

- moving records for free space management changes rid;
- may not be acceptable.

Page Formats: Variable Length Records

slot's format: $\langle \text{record offset, record length} \rangle$



✉ *Can move records on page without changing rid; so, attractive for fixed-length records too.*

Slotted page: a detailed view

slot's format: <record offset, record length>

