

Relational Query Optimization

SUN YAT-SEN UNIVERSITY

Review

■ 单个关系运算的求值方法

- 选择运算: **cheapest access path** (文件扫描、索引)
- 投影运算: 消除重复 (排序)
- 连接运算: **Nested Loops Join**、**Sort-Merge Join**

■ 相应的开销估算公式

- 外排序 $2N(1 + \lceil \log_{B-1} \lceil N/B \rceil \rceil)$
- 块嵌套循环连接算法 **Cost = [R] + [R]/N * [S]**
- 排序归并连接算法 **Cost: Sort R + Sort S + ([R]+[S])**

Review

■ Choice of single-table operations

- Depends on indexes, memory, stats,...

■ Joins

- Blocked nested loops:
 - simple, exploits extra memory
 - Indexed nested loops:
 - best if 1 rel small and one indexed
 - Sort/Merge Join
 - good with small amount of memory, bad with duplicates
 - Hash Join
 - fast (enough memory), bad with skewed data
-


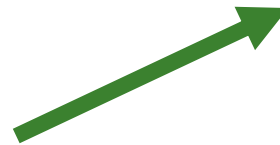
Query Optimization Overview

- Query can be converted to relational algebra
 - Relational Algebra converts to tree, joins form branches

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```



$\pi_{(sname)} \sigma_{(bid=100 \wedge rating > 5)} (Reserves \bowtie Sailors)$



- Each operator has implementation choices
- Operators can also be applied in different order!

Query Optimization Overview (cont.)

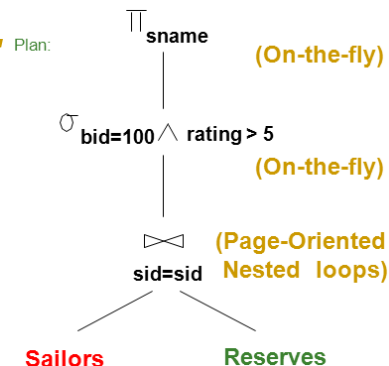
- Plan(执行计划, 查询求解计划): *Tree of Relation Algebra operations (and some others) with choice of algorithm for each operation.*

- Three main issues:

- For a given query, what plans are considered?
- How is the cost of a plan estimated?
- How do we “search” in the “plan space”?

- Ideally: Want to find best plan.

- Reality: Avoid worst plans!



Cost-based Query Sub-System

Queries

```
Select *  
From Blah B  
Where B.blah = blah
```

Query Parser

Query Optimizer

Plan
Generator

Plan Cost
Estimator


Usually there is a **heuristics**-based rewriting step before the cost-based steps.

Catalog Manager

Schema

Statistics

Query Executor



Schema for Examples

Sailors (sid: integer, sname: string, rating: integer, age: real)

Reserves (sid: integer, bid: integer, day: dates, rname: string)

■ Reserves:

- ❑ Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
- ❑ Assume there are 100 boats

[R]=1000, $p_R=100$
100 boats

■ Sailors:

- ❑ Each tuple is 50 bytes long, 80 tuples per page, 500 pages.
- ❑ Assume there are 10 different ratings

[S]=500, $p_S=80$.
10 ratings

■ Assume we have 5 pages in our buffer pool!

Motivating Example

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

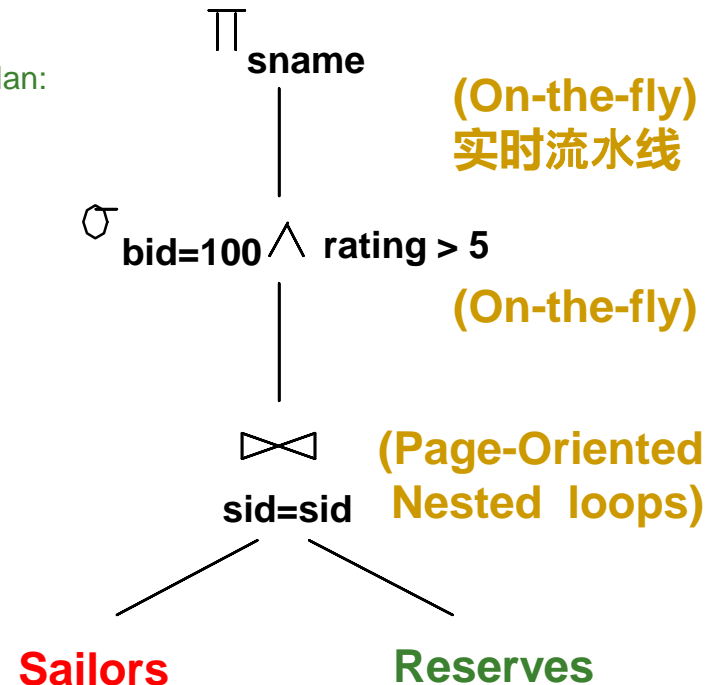
[R]=1000, $p_R=100$
100 boats
[S]=500, $p_S=80$.
10 ratings
5 buffer pages

Plan:

- Cost: $500 + 500 * 1000$ I/Os = 500,500 I/Os

$$\text{Cost} = [O] + [O] * [I]$$

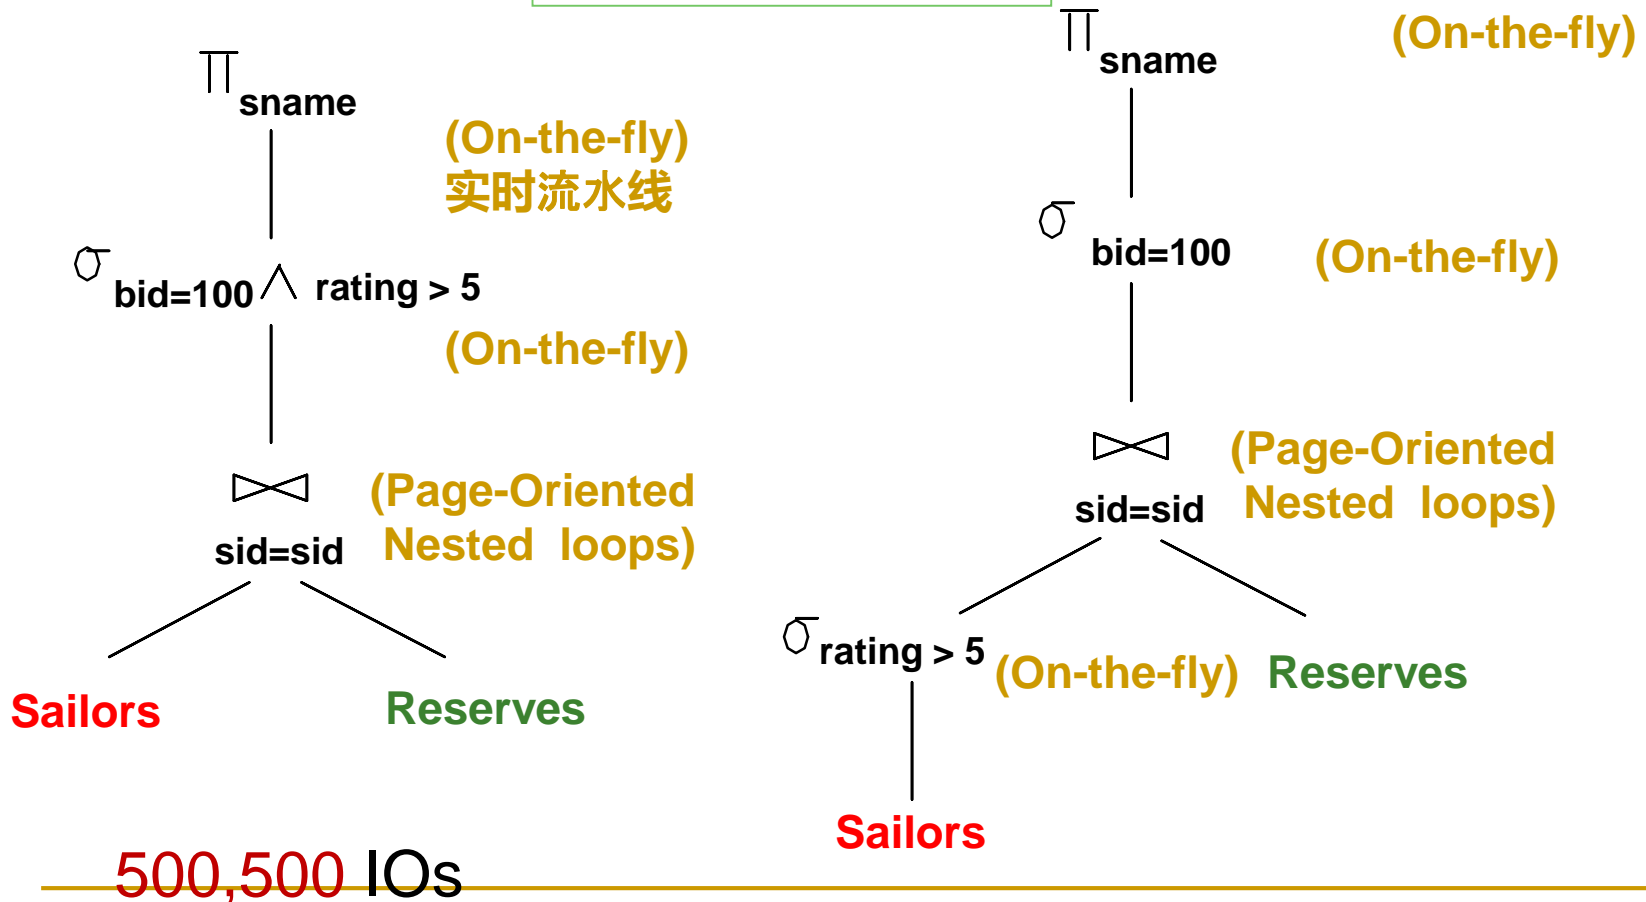
- By no means the worst plan!
- Misses several opportunities:
 - selections could be **pushed** down(下推)
 - no use made of indexes
- *Goal of optimization:* Find faster plans that compute the same answer.



Alternative Plans – Push Selects (No Indexes) 下推选择

[R]=1000, $p_R=100$
100 boats
[S]=500, $p_S=80$.
10 ratings
5 buffer pages

$$\text{Cost} = [O] + [O] * [I]$$

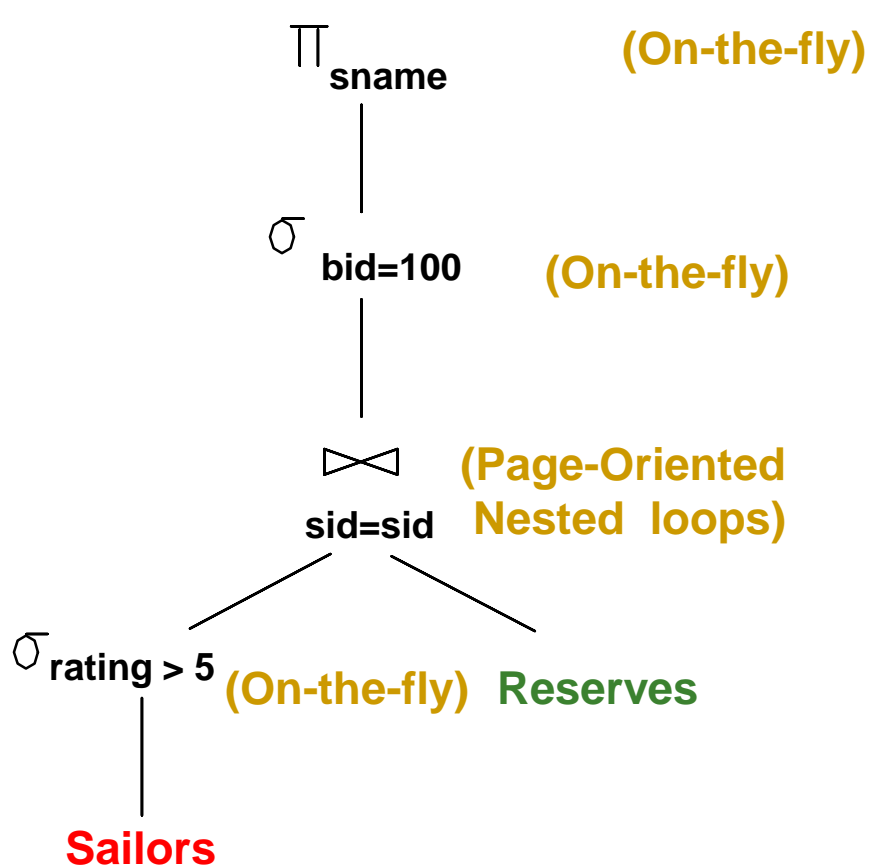


$$500 + 250 + 250 * 1000 = 250,500 \text{ IOs}$$

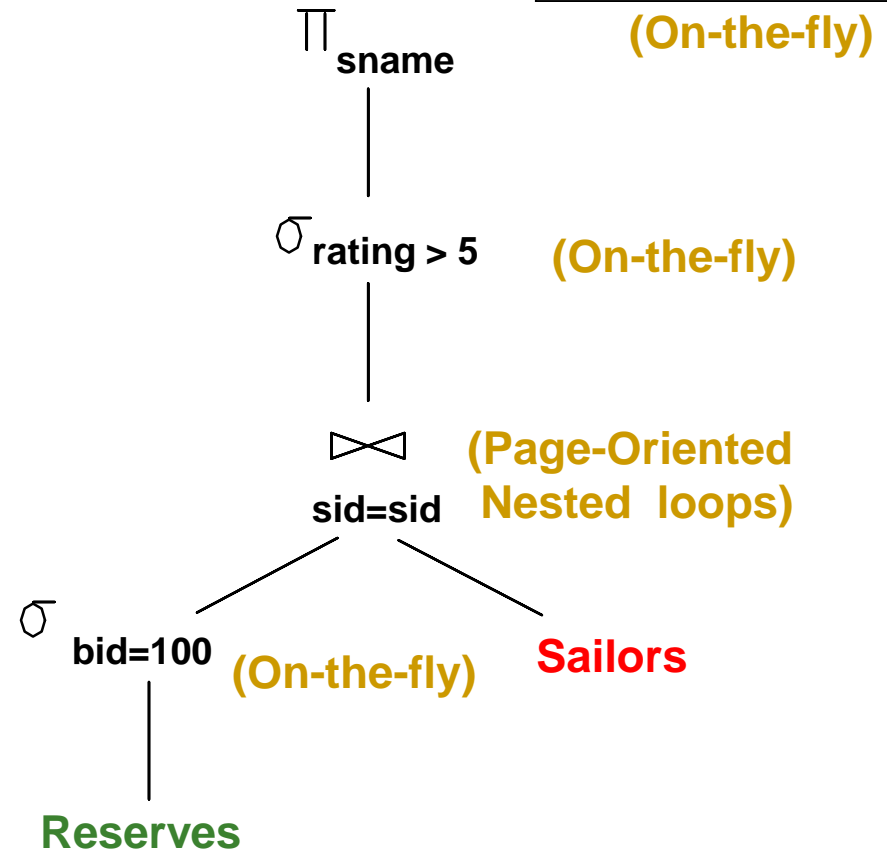
Alternative Plans – Push Selects (No Indexes)

$$\text{Cost} = [O] + [O] * [I]$$

[R]=1000, $p_R=100$
100 boats
[S]=500, $p_S=80$.
10 ratings
5 buffer pages



250,500 IOs

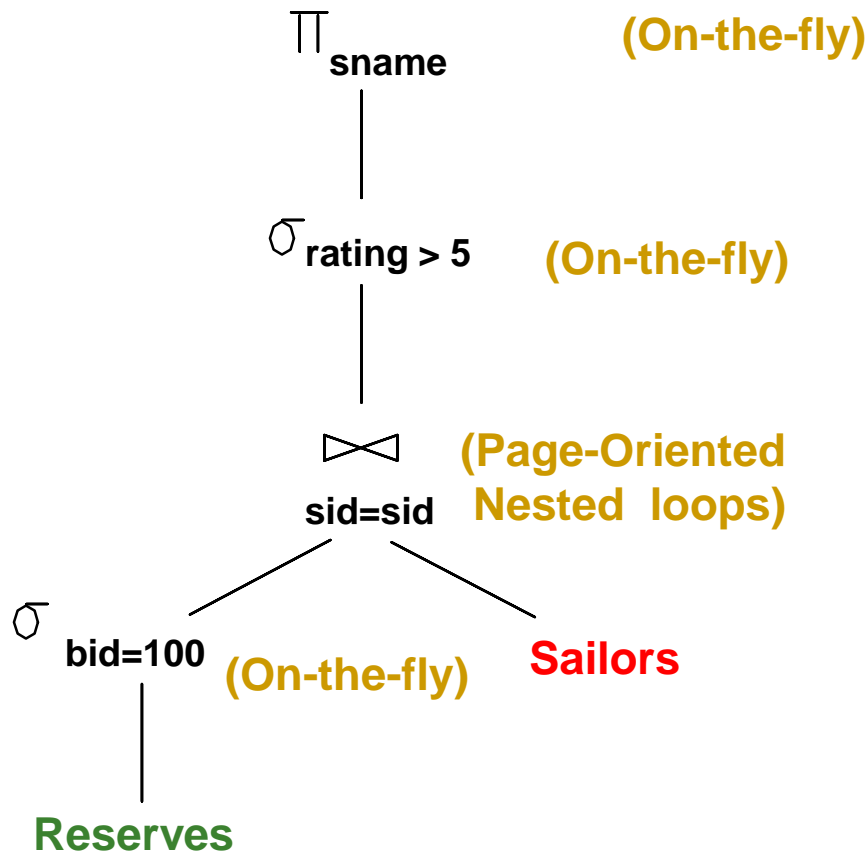


1000 + 10*500=6000 IOs

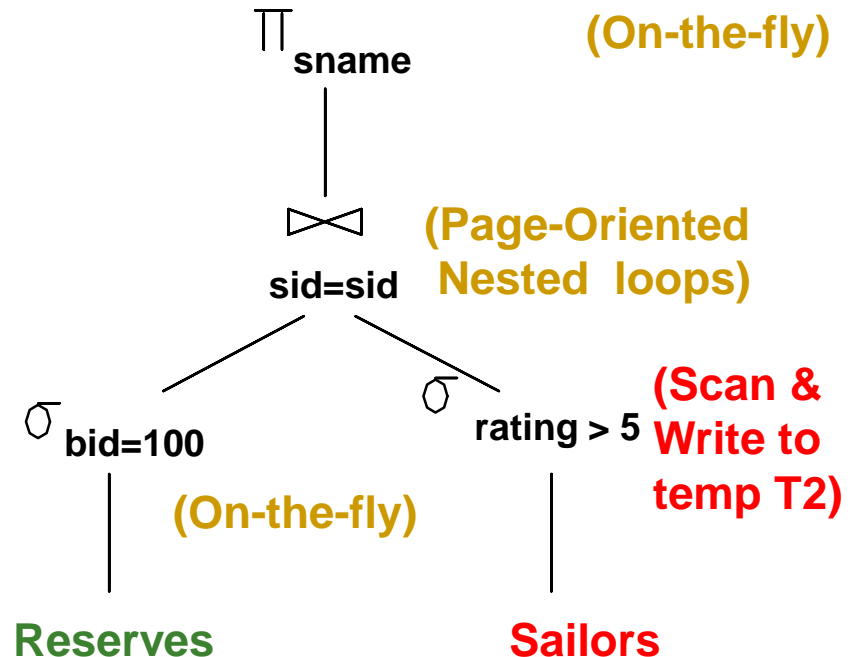
Alternative Plans – Push Selects (No Indexes)

$$\text{Cost} = [O] + [O] * [I]$$

[R]=1000, $p_R=100$
100 boats
[S]=500, $p_S=80$.
10 ratings
5 buffer pages



6000 IOs

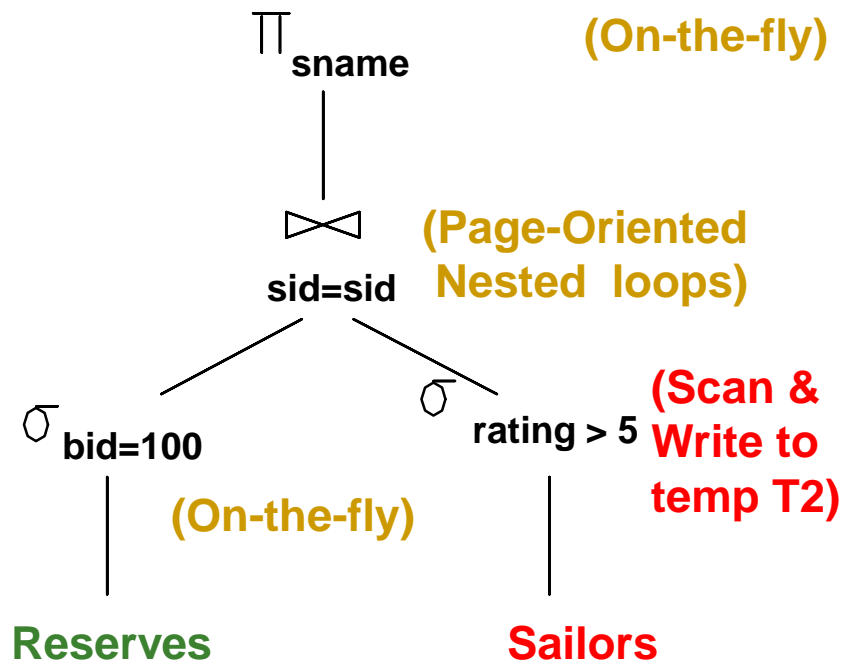


$1000 + 500 + 250 + (10 * 250) = 4250$ IOs

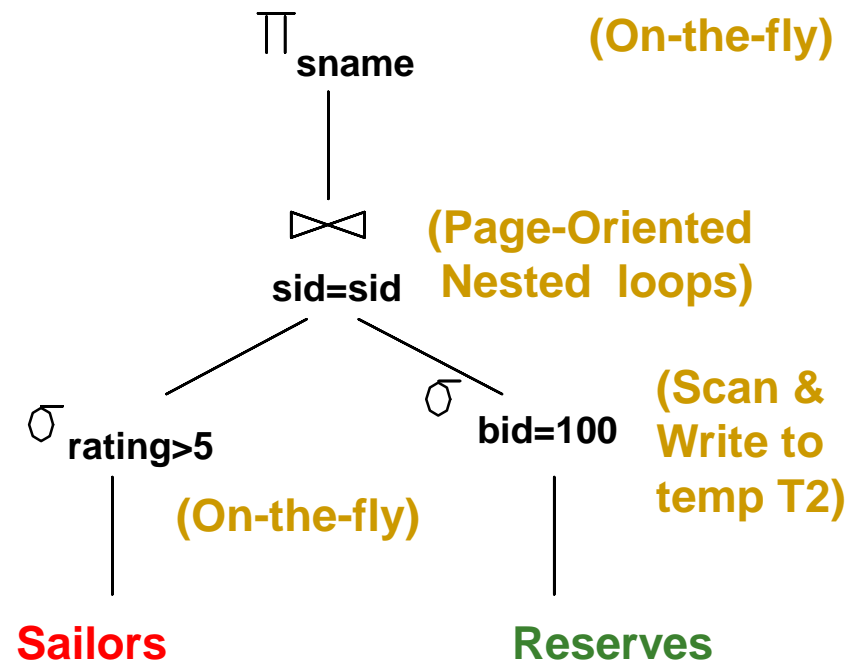
Alternative Plans – Push Selects (No Indexes)

$$\text{Cost} = [O] + [O] * [I]$$

[R]=1000, $p_R=100$
100 boats
[S]=500, $p_S=80$.
10 ratings
5 buffer pages

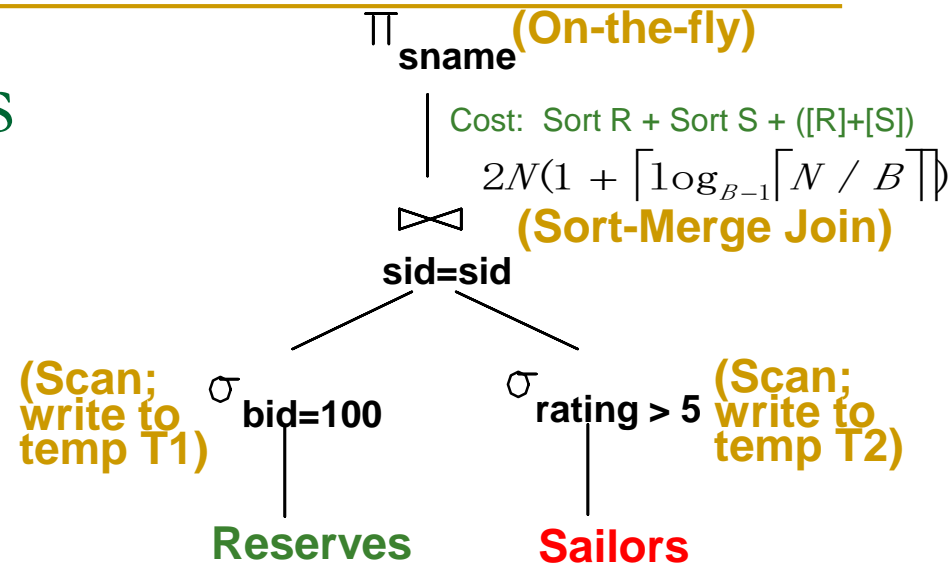


4250 IOs



$500 + 1000 + 10 + (250 * 10) = 4010$ IOs

More Alternative Plans (No Indexes)



Sort Merge Join

With 5 buffers, cost of plan:

- ❑ Scan Reserves (1000) + write temp T1 (10 pages) = 1010.
- ❑ Scan Sailors (500) + write temp T2 (250 pages) = 750.
- ❑ Sort T1 ($2 \cdot 2 \cdot 10$) + sort T2 ($2 \cdot 4 \cdot 250$) + merge (10+250) = 2300
- ❑ Total: 4060 page I/Os.

$$\text{Cost} = [O] + [O]/N * [I]$$

If use BNL join, join = $10 + 4 \cdot 250$, total cost = 2770.

Can also 'push' projections, but must be careful!

- ❑ T1 has only *sid*, T2 only *sid*, *sname*:
- ❑ T1 fits in 3 pgs, cost of BNL under 250 pgs, total < 2000.

[R]=1000, $p_R=100$
100 boats
[S]=500, $p_S=80$.
10 ratings
5 buffer pages

Summing up

- There are *lots* of plans
 - Even for a relatively simple query
 - People tend to think they can pick good ones by hand
 - MapReduce is based on that assumption
 - Not so clear that's true!
 - Machines are better at enumerating options than people
 - But we will see soon how optimizers make simplifying assumptions
-

What is Needed for Optimization?

- A closed set of operators
 - Relational ops (table in, table out)
 - Encapsulation (e.g. based on iterators)
- Plan space
 - Based on relational equivalences, different implementations

- Cost Estimation, based on

- Cost formulas
- Size estimation, in turn based on
 - Catalog information on base tables
 - Selectivity (Reduction Factor) estimation

$$\text{Cost} = [R] + [R]/N * [S] \\ 2N(1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil)$$

[R]=1000, $p_R=100$
100 boats
[S]=500, $p_S=80$.
10 ratings
5 buffer pages

- A search algorithm: To sift through the plan space and find lowest cost option!

Query Optimization

- Will focus on “System R” (Selinger) style optimizers

Access Path Selection
in a Relational Database Management System

P. Griffiths Selinger
M. M. Astrahan
D. D. Chamberlin
R. A. Lorie
T. G. Price

IBM Research Division, San Jose, California 95193



ABSTRACT: In a high level query and data manipulation language such as SQL, requests are stated non-procedurally, without reference to access paths. This paper describes how System R chooses access paths

retrieval. Nor does a user specify in what order joins are to be performed. The System R optimizer chooses both join order and an access path for each table in the SQL statement. Of the many possible

Highlights of System R Optimizer

■ Impact:

- ❑ Most widely used currently; works well for 10-15 joins.

■ Cost estimation:

- ❑ Very inexact, but works OK in practice.
 - ❑ Statistics in system catalogs used to estimate cost of operations and result sizes.
 - ❑ Considers combination of CPU and I/O costs.
-

Highlights of System R Optimizer (Contd)

- **Plan Space:** Too large, must be pruned.
 - Many plans share common, “overpriced” subtrees
 - ignore them all!
 - In some implementations, only the space of *left-deep plans* (左深计划) is considered.
 - Cartesian products avoided in some implementations.

Query Blocks: Units of Optimization

查询块

- Break query into *query blocks*
- Optimized one block at a time
- Uncorrelated nested blocks computed once
- Correlated nested blocks like function calls
 - But sometimes can be “decorrelated”
 - Beyond the scope of introductory course!


```
SELECT S.sname
FROM Sailors S
WHERE S.age IN
      (SELECT MAX (S2.age)
       FROM Sailors S2
       GROUP BY S2.rating)
```

Outer block

Nested block

❖ For each block, the plans considered are:

- All available *access methods*, for each relation in FROM clause.
- All *left-deep join trees*
 - right branch always a base table
 - consider all join *orders* and join *methods*



Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)
Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

■ Reserves:

- Each tuple is 40 bytes long, 100 tuples per page, 1000 pages. 100 distinct bids.

[R]=1000, $p_R=100$
100 boats
[S]=500, $p_S=80$.
10 ratings
40,000 sids

■ Sailors:

- Each tuple is 50 bytes long, 80 tuples per page, 500 pages. 10 ratings, 40,000 sids.

Translating SQL to Relational Algebra

```
SELECT S.sid, MIN (R.day)
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = "red"
GROUP BY S.sid
HAVING COUNT (*) >= 2
```

For each sailor with at least two reservations for red boats, find the sailor id and the earliest date on which the sailor has a reservation for a red boat.

Translating SQL to Relational Algebra


```
SELECT S.sid, MIN (R.day)
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = "red"
GROUP BY S.sid
HAVING COUNT (*) >= 2
```

π S.sid, MIN(R.day)
(HAVING COUNT(*)>2 (
GROUP BY S.Sid (
 $\sigma_{B.color = \text{"red"}}$ (
Sailors \bowtie Reserves \bowtie Boats))))

Relational Algebra Equivalences

- Allow us to choose different join orders and to 'push' selections and projections ahead of joins.
- Selections:
 - $\sigma_{c_1 \wedge \dots \wedge c_n}(R) \equiv \sigma_{c_1}(\dots(\sigma_{c_n}(R))\dots)$ (*cascade-级联*)
 - $\sigma_{c_1}(\sigma_{c_2}(R)) \equiv \sigma_{c_2}(\sigma_{c_1}(R))$ (*commute-交换律*)
- Projections:
 - $\pi_{a_1}(R) \equiv \pi_{a_1}(\dots(\pi_{a_1, \dots, a_n}(R))\dots)$ (*cascade*)
- Cartesian Product
 - $(R \times S) \times T \equiv R \times (S \times T)$ (*associative-结合律*)
 - $R \times S \equiv S \times R$ (*commutative*)
 - *This means we can do joins in any order.*

More Equivalences



■ Eager projection


- ❑ Can cascade and “push” some projections thru selection
- ❑ Can cascade and “push” some projections below one side of a join
- ❑ Rule of thumb: can project anything not needed “downstream”

■ Selection on a cross-product is equivalent to a join.


- ❑ If selection is comparing attributes from each side

■ A selection on attributes of R commutes with $R \bowtie S$.

- ❑ i.e., $\sigma(R \bowtie S) \equiv \sigma(R) \bowtie S$
- ❑ but only if the selection doesn't refer to S!



Cost Estimation



- For each plan considered, must estimate total cost:
 - Must estimate *cost* of each operation in plan tree.
 - We've already discussed this for various operators
 - sequential scan, index scan, joins, etc.
 - Depends on input cardinalities(基数).
- Cost = [R] + [R]/N * [S]**
 $2N(1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil)$
- Must estimate *size of result* for each operation in tree!
 - Use information about the input relations.
- Q: Is “cost” the same as estimated “run time”?

Statistics and Catalogs

- Need information on relations and indexes involved.
- *Catalogs* typically contain at least:

Statistic	Meaning
NTuples	# of tuples in a table (cardinality)
NPages	# of disk pages in a table
Low/High	min/max value in a column
Nkeys	# of distinct values in a column
IHeight	the height of an index
INPages	# of disk pages in an index

- Catalogs updated periodically.
- Modern systems do more
 - keep more detailed information on data values, e.g., histograms

Size Estimation and Selectivity

```
SELECT  attribute list  
FROM    relation list  
WHERE   term1 AND ... AND termk
```

- Max output cardinality = product of input cardinalities
- *Selectivity (sel)* associated with each *term*
 - Book calls selectivity “Reduction Factor” (RF)
 - $|\text{output}| / |\text{input}|$
 - reflects the impact of the *term* in reducing result size.

Result cardinality = Max # tuples * $\prod \text{sel}_i$

Reduction Factor Estimation

```
SELECT  attribute list
FROM    relation list
WHERE   term1 AND ...
```

■ *Result cardinality* = Max # tuples * product of all RF's.

■ Term *col=value* (given Nkeys(I) on col)
RF = $1/NKeys(I)$

■ Term *col1=col2* (handy for joins too...)
RF = $1/MAX(NKeys(I1), NKeys(I2))$

■ Term *col>value*
RF = $(High(I)-value)/(High(I)-Low(I))$

Implicit assumptions: **values are uniformly distributed and**
terms are independent!


■ *Note, if missing the needed stats, assume 1/10!!!*

Reduction Factors & Histograms

- For better estimation, use a histogram


equiwidth

No. of Values	2	3	3	1	8	2	1
Value	0-.99	1-1.99	2-2.99	3-3.99	4-4.99	5-5.99	6-6.99




equidepth

No. of Values	2	3	3	3	3	2	4
Value	0-.99	1-1.99	2-2.99	3-4.05	4.06-4.67	4.68-4.99	5-6.99



Think Through Estimation for Joins

- Term $col1=col2$
 - $RF = 1/MAX(NKeys(I1), NKeys(I2))$
- Q: Given $R \text{ join } S$, what is result cardinality?
Two cases:
 1. Key for R , Foreign Key for S
 - A common case, treat it specially! $RF = 1/|R|$
 2. join on non-key $\{A\}$
 - For each $r \in R$, $NTuples(S)/NKeys(A,S)$ result tuples
so... $(NTuples(\mathbf{R}) * NTuples(\mathbf{S})) / NKeys(A, \mathbf{S})$
 - For each $s \in S$, $NTuples(R)/NKeys(A, \mathbf{R})$
so... $(NTuples(\mathbf{S}) * NTuples(\mathbf{R})) / NKeys(A, \mathbf{R})$
- If these two estimates differ, take the lower one!
 - Q: Why?


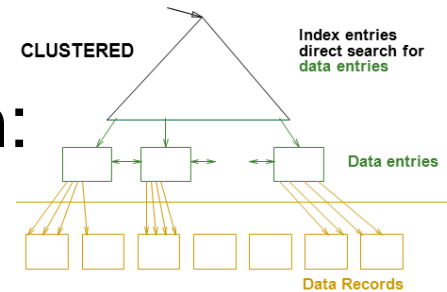


Enumeration of Alternative Plans

- There are two main cases:
 - Single-relation plans (base case)
 - Multiple-relation plans (induction)
- Single-table queries include selects, projects, and grouping/aggregate operations:
 - Consider each available access path (file scan / index)
 - Choose the one with the least estimated cost
 - Selection/Projection done on the fly
 - Result pipelined into grouping/aggregation

Cost Estimates for Single-Relation Plans

- Index I on primary key matches selection:
 - Cost is $Height(I)+1$ for a B+ tree.
- Clustered index I matching one or more selects:
 - $(NPages(I)+NPages(R)) * \text{product of RF's of matching selects}$.
- Non-clustered index I matching one or more selects:
 - $(NPages(I)+NTuples(R)) * \text{product of RF's of matching selects}$.
- Sequential scan of file:
 - $NPages(R)$.



■ Recall: Must also charge for duplicate elimination if required

Example

```
SELECT S.sid
FROM Sailors S
WHERE S.rating=8
```

[R]=1000, $p_R=100$
100 boats
[S]=500, $p_S=80$.
10 ratings
5 buffer pages

■ If we have an **index on *rating***:


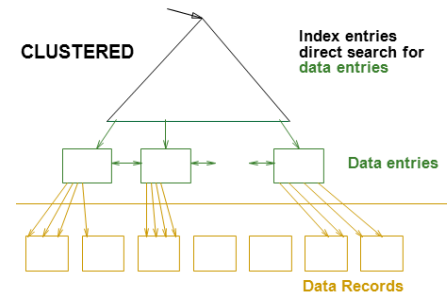
- ❑ Cardinality = $NTuples(S) * (1/NKeys(I)) = 40000 \text{ tuples} * (1/10)$
- ❑ **Clustered index:** $(NPages(I)+NPages(S)) * (1/NKeys(I))$
 $= (50+500) * (1/10) = 55 \text{ pages are retrieved. (This is the **cost**.)}$
- ❑ **Unclustered index:** $(NPages(I)+NTuples(S)) * (1/NKeys(I))$
 $= (50+40000) * (1/10) = 4005 \text{ pages are retrieved.}$

■ If we have an **index on *sid***:

- ❑ Would have to retrieve all tuples/pages.
 - With a **clustered** index, the **cost** is $50+500$,
 - with **unclustered** index, $50+40000$.


■ Doing a **file scan**:

- ❑ We retrieve all file pages (**500**).




Queries Over Multiple Relations

- A System R heuristic:
only left-deep join trees considered.



- ❑ Restricts the search space
- ❑ Left-deep trees allow us to generate all *fully pipelined plans*.
 - Intermediate results not written to temporary files.
 - Not all left-deep trees are fully pipelined (e.g., SM join).

Enumeration of Left-Deep Plans



- Left-deep plans differ in
 - the order of relations
 - the access method for each relation
 - the join method for each join.
- Enumerated using N passes (if N relations joined):
 - **Pass 1:** Find best 1-relation plan for each relation.
 - **Pass i:** Find best way to join result of an (i - 1)-relation plan (as outer) to the i'th relation. (i between 2 and N.)
- For each subset of relations, retain only:
 - Cheapest plan overall, plus
 - Cheapest plan for each *interesting order* of the tuples.

The Dynamic Programming Table


[R]=1000, $p_R=100$
100 boats
[S]=500, $p_S=80$.
10 ratings
5 buffer pages

Subset of tables in FROM clause	Interesting-order columns	Best plan	Cost
{R, S}	<none>	hashjoin(R,S)	4500
{R, S}	<R.a, S.b>	sortmerge(R,S)	7500

A Note on “Interesting Orders”

- An intermediate result has an “interesting order” if it is sorted by any of:
 - ORDER BY attributes
 - GROUP BY attributes
 - Join attributes of *yet-to-be-added* (downstream) joins

Enumeration of Plans (Contd.)



- Match an $i - 1$ way plan with another table *only if*
 - a) there is a join condition between them, *or*
 - b) all predicates in WHERE have been used up.
 - i.e., avoid Cartesian products if possible.
- ORDER BY, GROUP BY, aggregates etc. handled as a final step
 - via 'interestingly ordered' plan if chosen (free!)
 - or via an additional sort/hash operator
- Despite pruning, this is exponential in #tables.

Example

Sailors:

Hash, B+ on *sid*

Reserves:

Clustered B+ tree on *bid*

B+ on *sid*

Boats

B+ on *color*

Sid, COUNT(*) AS numbes

GROUPBY_{sid}

sid=sid

bid=bid

Sailors

Reserves

Color=red

Boats

Select S.sid, COUNT(*) AS number
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
AND B.color = "red"
GROUP BY S.sid

- **Pass1: Best plan(s) for accessing each relation**
 - Reserves, Sailors: File Scan
 - Q: What about Clustered B+ on Reserves.bid???
 - Boats: B+ tree on color

Pass 1

- Find best plan for each relation in isolation:
 - Reserves, Sailors: File Scan
 - Boats: B+ tree on color


FROM Sailors S, Reserves R, Boats B

Pass 2

- For each plan in pass 1, generate plans joining another relation as the inner, using all join methods (and **matching inner access methods**)

- File Scan Reserves (outer) with Boats (inner)
- File Scan Reserves (outer) with Sailors (inner)
- File Scan Sailors (outer) with Boats (inner)
- File Scan Sailors (outer) with Reserves (inner)
- Boats Btree on color with Sailors (inner)
- Boats Btree on color with Reserves (inner)

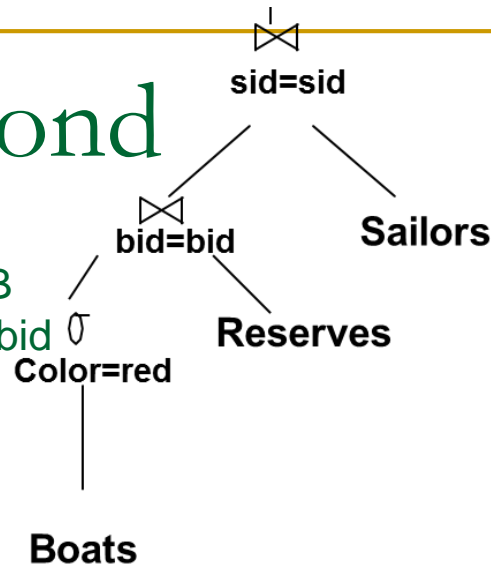
- Retain cheapest plan for each (pair of relations, order)



FROM Sailors S, Reserves R, Boats B

Pass 3 and beyond

```
Select S.sid, COUNT(*) AS number
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
      AND B.color = "red"
GROUP BY S.sid
```



Sailors:

Hash, B+ on *sid*

Reserves:

Clustered B+ tree on *bid*

B+ on *sid*

Boats

B+ on *color*

- Using **Pass 2 plans** as outer relations, generate plans for the next join
 - E.g. **Boats B+-tree on color with Reserves (bid) (sortmerge)**
inner Sailors (B-tree *sid*) sort-merge
- Then, add cost for groupby/aggregate:
 - This is the cost to sort the result by *sid*, *unless it has already been sorted by a previous operator*.
- Then, choose the cheapest plan

FROM Sailors S, Reserves R, Boats B

Summary

- Optimization is the reason for the lasting power of the relational system
- But it is primitive in some ways
- New areas: many!
 - Smarter summary statistics (fancy histograms and “sketches”)
 - Auto-tuning statistics,
 - Adaptive runtime re-optimization (e.g. *eddies*),
 - Multi-query optimization,
 - And parallel scheduling issues, etc.

Summary

■ 要求:

- 理解查询树、执行计划树、关系代数等价规则、左深计划树和完全流水线计划树等概念
- 深刻理解选择条件的选择性/缩减因子，进而能够估算结果集大小
- 能够估算执行计划的开销
- 理解左深计划的遍历算法
 - 若只有1、2关系，则能够找出最优计划