

Deep Feedforward Networks

Hankz Hankui Zhuo

May 10, 2019

<http://xplan-lab.org>

What's DFN

- Information **flows through** the function being evaluated from **x**, through the intermediate computations used to define **f**, and finally to the output **y**
 - **No feedback connections**
- Extended to include feedback connections
 - **Recurrent** Neural Networks
- Also called **feedforward neural networks**
- Or **multilayer perceptions (MLPs)**

Structure

- Chain structures:
 - three layers: $f(x) = f^3(f^2(f^1(x)))$
 - “three” is the **depth** of the model
 - **Output layer**: the desired output is specified in the training data
 - **Hidden layers**: the desired output is not specified in the training data
 - **Width of the model**: the dimensionality of hidden layers

Mapping of input

- Linear model of input x
 - $y = w^T x$
- Nonlinear model
 - Introducing mapping ϕ
 - $y = w^T \phi(x)$
- How to choose ϕ ?
 - 1. use a very generic ϕ :
 - kernel machines, e.g., RBF kernel
 - If $\phi(x)$ is of high enough dimension, we can always have enough capacity to fit the training set
 - 2. manually engineer ϕ


Mapping of input

- 3. The strategy of deep learning is to learn ϕ
 - $y=f(x;\theta,w)=\phi(x; \theta)^T w$
 - This approach can capture the benefit of the first approach by being highly generic
 - we do so by using a very broad family $\phi(x;\theta)$.
 - Can also capture the benefit of the second approach
 - Human practitioners can encode their knowledge to help generalization by designing families $\phi(x; \theta)$ that they expect will perform well.

Example: Learning XOR

- XOR is an operation on two binary values
- When exactly one of these binary values is equal to 1, the XOR function returns 1
- Training data with four points
 - $X = \{[0,0], [0,1], [1,0], [1,1]\}$
- MSE (mean squared error) loss function:
 - $J(\theta) = 1/4 \sum_x (f(x) - f(x; \theta))^2$
- Choose the form of our model $f(x; \theta)$
 - Linear model: $f(x; w, b) = x^T w + b$
- We got $w=0$ and $b=1/2$, i.e., The linear model simply outputs 0.5 everywhere


Why?



- When $x_1 = 0$, the model's output must increase as x_2 increases.
- When $x_1 = 1$, the model's output must decrease as x_2 increases.
- A linear model must apply a fixed coefficient w_2 to x_2 .
- The linear model therefore cannot use the value of x_1 to change the coefficient on x_2 and cannot solve this problem.


A simple feedforward network

- One hidden layer with two hidden units
- Hidden layer:
 - $h = f^1(x; W, c)$
- Output layer:
 - $y = f^2(h; w, b)$
- Complete model:
 - $f(x; W, c, w, b) = f^2(f^1(x))$
- If f^1 use linear model, then
 - $h = f^1(x) = W^T x$, $f^2(h) = h^T w$,
 - Then we have $y = f(x) = w^T W^T x$
 - Or just $f(x) = x^T w'$, where $w' = Ww$
 - We thus cannot solve the problem
- We need a nonlinear function
- Use a fixed nonlinear function called activation function
 - $h = g(W^T x + c)$
- g is typically chosen to be a function that is applied element-wise
 - $h_i = g(x^T W_{:,i} + c_i)$



Rectified Linear Unit (ReLU)

- The activation function g is defined by
 - $g(z) = \max\{0, z\}$
- i.e., the complete network is:
 - $f(x; W, c, w, b) = w^T \max\{0, W^T x + c\} + b$



A solution to XOR

- A solution as show below:

$$W = \begin{bmatrix} 1 & 1 \\ 1 & 1 \end{bmatrix}, \quad w = \begin{bmatrix} 1 \\ -2 \end{bmatrix}$$

$$c = \begin{bmatrix} 0 \\ -1 \end{bmatrix}, \quad b = 0$$

- We can walk through the way the model processes based on $f(x;W,c,w,b)=w^T \max\{0,W^T x+c\}+b$

$$X = \begin{bmatrix} 0 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 1 \end{bmatrix} \quad XW = \begin{bmatrix} 0 & 0 \\ 1 & 1 \\ 1 & 1 \\ 2 & 2 \end{bmatrix} \quad \begin{bmatrix} 0 & -1 \\ 1 & 0 \\ 1 & 0 \\ 2 & 1 \end{bmatrix} \quad \begin{bmatrix} 0 & 0 \\ 1 & 0 \\ 1 & 0 \\ 2 & 1 \end{bmatrix} \quad \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \end{bmatrix}$$

add c apply ReLU multiply w

Gradient-Based Learning

- difference between the linear models and neural networks is
 - the nonlinearity of a neural network causes most interesting loss functions to become non-convex
 - This means that neural networks are usually trained by using iterative, gradient-based optimizers that merely drive the cost function to a very low value, rather than the linear equation solvers used to train linear regression models or the convex optimization algorithms with global convergence guarantees
 - Stochastic gradient descent applied to non-convex loss functions has **no such convergence guarantee**, and is sensitive to the values of the initial parameters.
 - For feedforward neural networks, it is important to initialize all weights to small random values.

Gradient-Based Learning

- The biases may be initialized to zero or to small positive values.
- For the moment, it suffices to understand that the training algorithm is almost always based on using the gradient to descend the cost function in one way or another.
- The specific algorithms are improvements and refinements on the ideas of gradient descent
- Computing the gradient is slightly more complicated for a neural network, but can still be done efficiently and exactly

Cost Function

- Cross-entropy between training data and the model distribution

$$J(\boldsymbol{\theta}) = -\mathbb{E}_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{\text{data}}} \log p_{\text{model}}(\mathbf{y} \mid \mathbf{x})$$

- The specific form of the cost function changes from model to model, depending on the specific form of $\log p_{\text{model}}$
- For $p_{\text{model}}(\mathbf{y} \mid \mathbf{x}) = \mathcal{N}(\mathbf{y}; f(\mathbf{x}; \boldsymbol{\theta}), \mathbf{I})$
- We have

$$J(\theta) = \frac{1}{2} \mathbb{E}_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{\text{data}}} \|\mathbf{y} - f(\mathbf{x}; \boldsymbol{\theta})\|^2 + \text{const}$$

Output Units

- suppose that the feedforward network provides a set of hidden features defined by
 - $h = f(x; \theta)$
- Linear Units for Gaussian Output Distributions

$$p(\mathbf{y} \mid \mathbf{x}) = \mathcal{N}(\mathbf{y}; \hat{\mathbf{y}}, \mathbf{I})$$

- Maximizing the log-likelihood is then equivalent to minimizing the mean squared error

Output Units (continued)

- **Sigmoid Units** for Bernoulli Output Distributions
- Suppose we were to use a linear unit, and threshold its value to obtain a valid probability

- A sigmoid $P(y = 1 \mid \mathbf{x}) = \max \left\{ 0, \min \left\{ 1, \mathbf{w}^\top \mathbf{h} + b \right\} \right\}$

$$\hat{y} = \sigma \left(\mathbf{w}^\top \mathbf{h} + b \right)$$

- logistic sigmoid

$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$

Output Units (continued)

- The sigmoid output unit as having two components
- First, it uses a linear layer to compute $z = w^T h + b$
- Next, it uses the sigmoid activation function to convert z into a probability
- We omit the dependence on x for the moment to discuss how to define a probability distribution over y using the value z

Output Units (continued)

- omit the dependence on x for the moment to discuss how to define a probability distribution over y using the value z
- The sigmoid can be motivated by constructing an unnormalized probability distribution $\tilde{P}(y)$
- We then normalize to see that this yields a Bernoulli distribution controlled by a sigmoidal transformation of z
- The loss function for maximum likelihood learning of a Bernoulli parametrized by a sigmoid is

$$\begin{aligned} J(\theta) &= -\log P(y \mid \mathbf{x}) \\ &= -\log \sigma((2y - 1)z) \\ &= \zeta((1 - 2y)z). \end{aligned}$$

$$\log \tilde{P}(y) = yz$$

$$\tilde{P}(y) = \exp(yz)$$

$$P(y) = \frac{\exp(yz)}{\sum_{y'=0}^1 \exp(y'z)}$$

$$P(y) = \sigma((2y - 1)z).$$

To be continued!