

Chapter 16 Meta Learning

Introduction

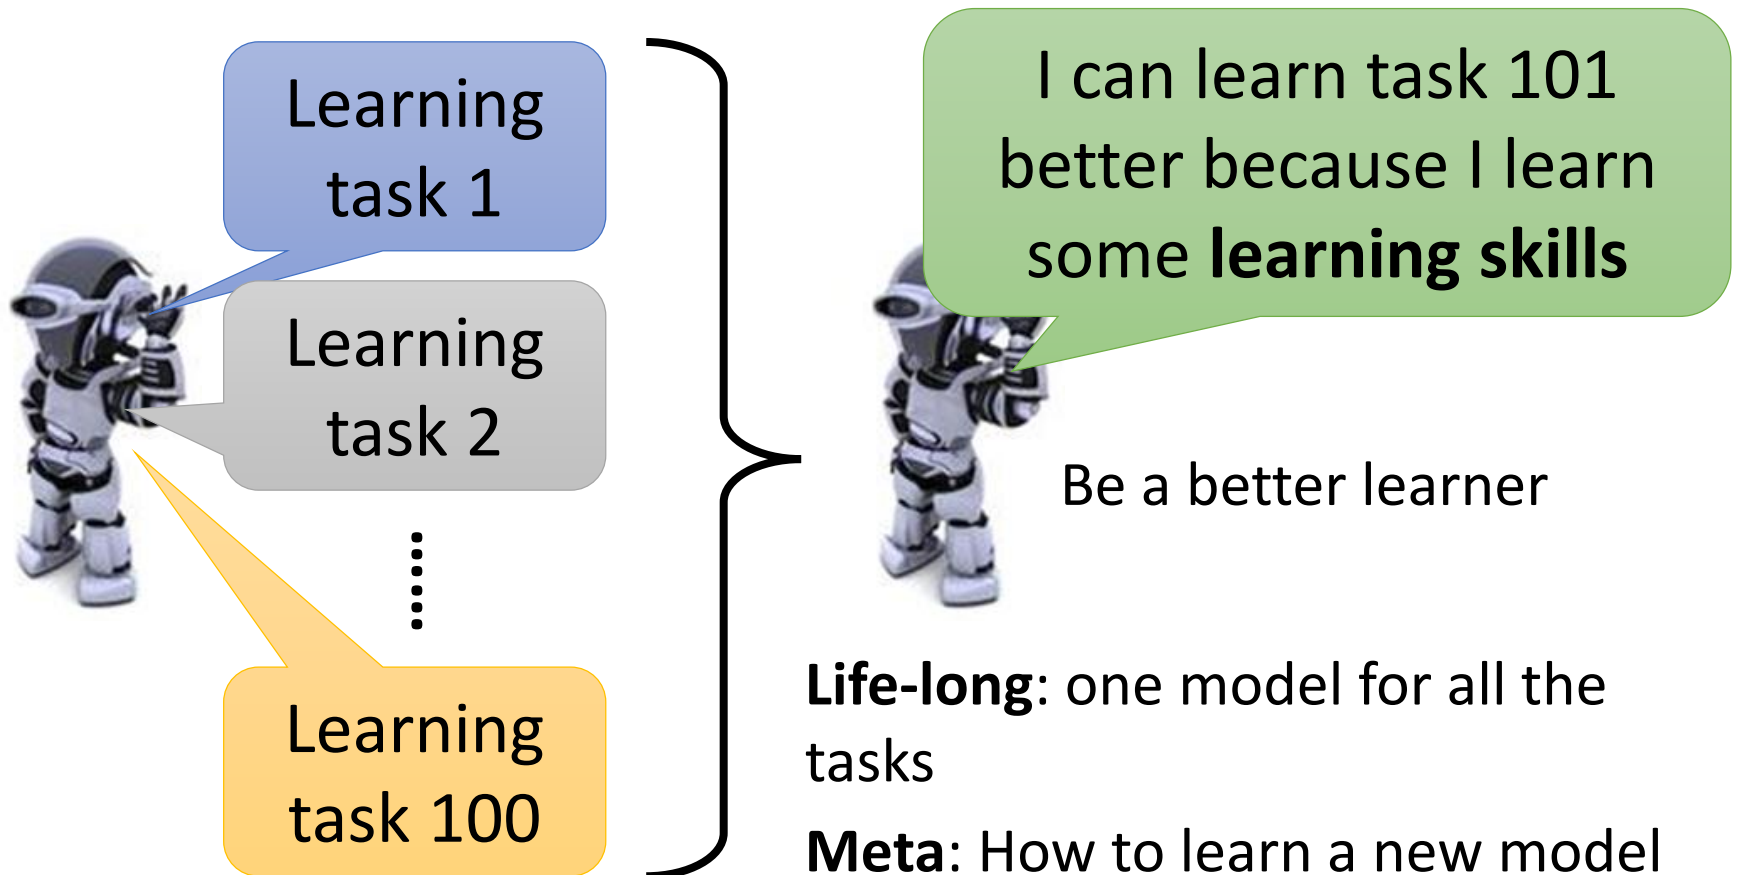
Task 1: speech recognition

Task 2: image recognition

⋮

Task 100: text classification

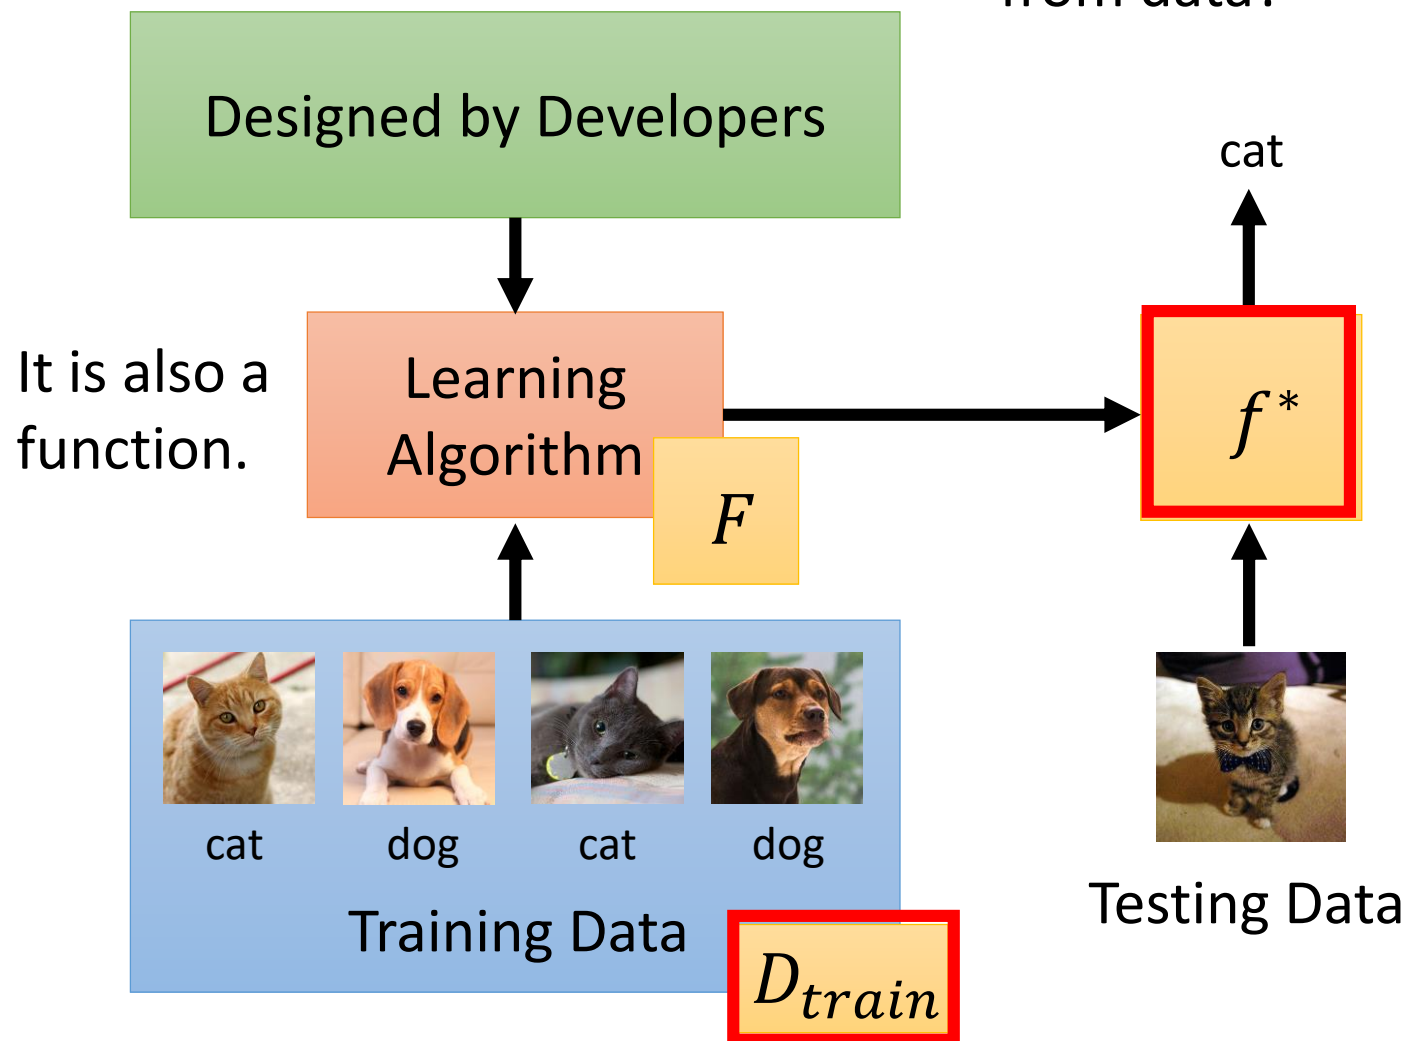
- Meta learning = Learn to learn



Meta Learning

$$f^* = F(D_{train})$$

Can machine find F
from data?



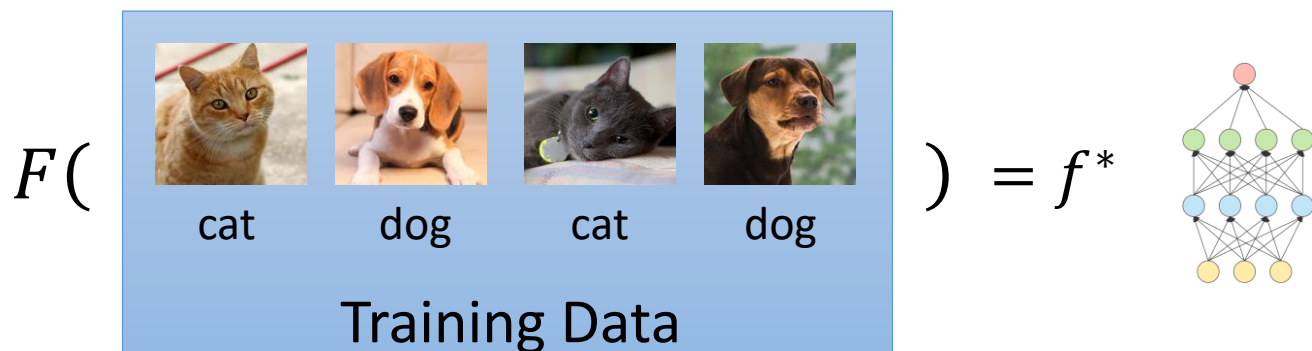
Meta Learning

Machine Learning \approx 根据资料找一个函数 f 的能力



Meta Learning

\approx 根据资料找一个找 一个函数 f 的函数 F 的能力



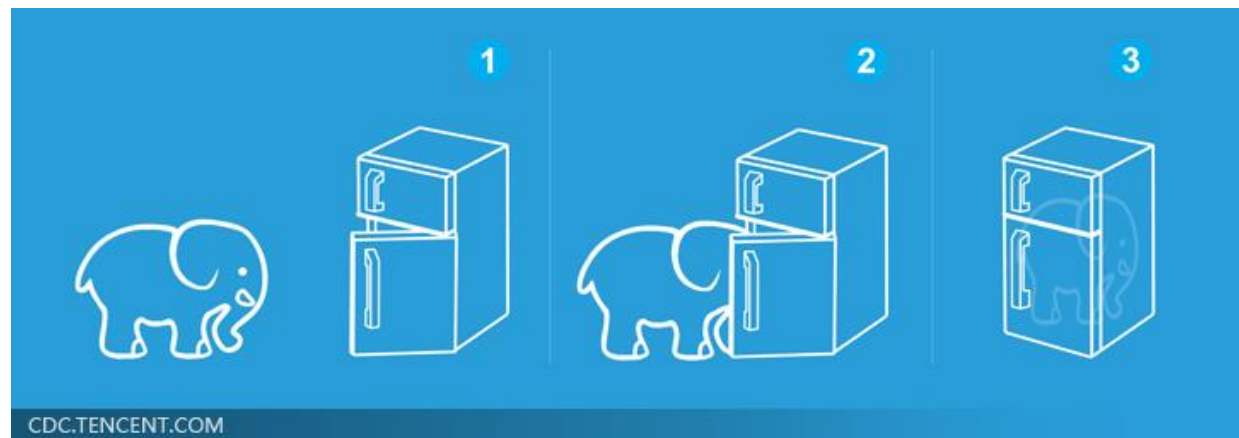
~~Machine~~ Learning is Simple

Meta



Function f \longrightarrow Learning algorithm F

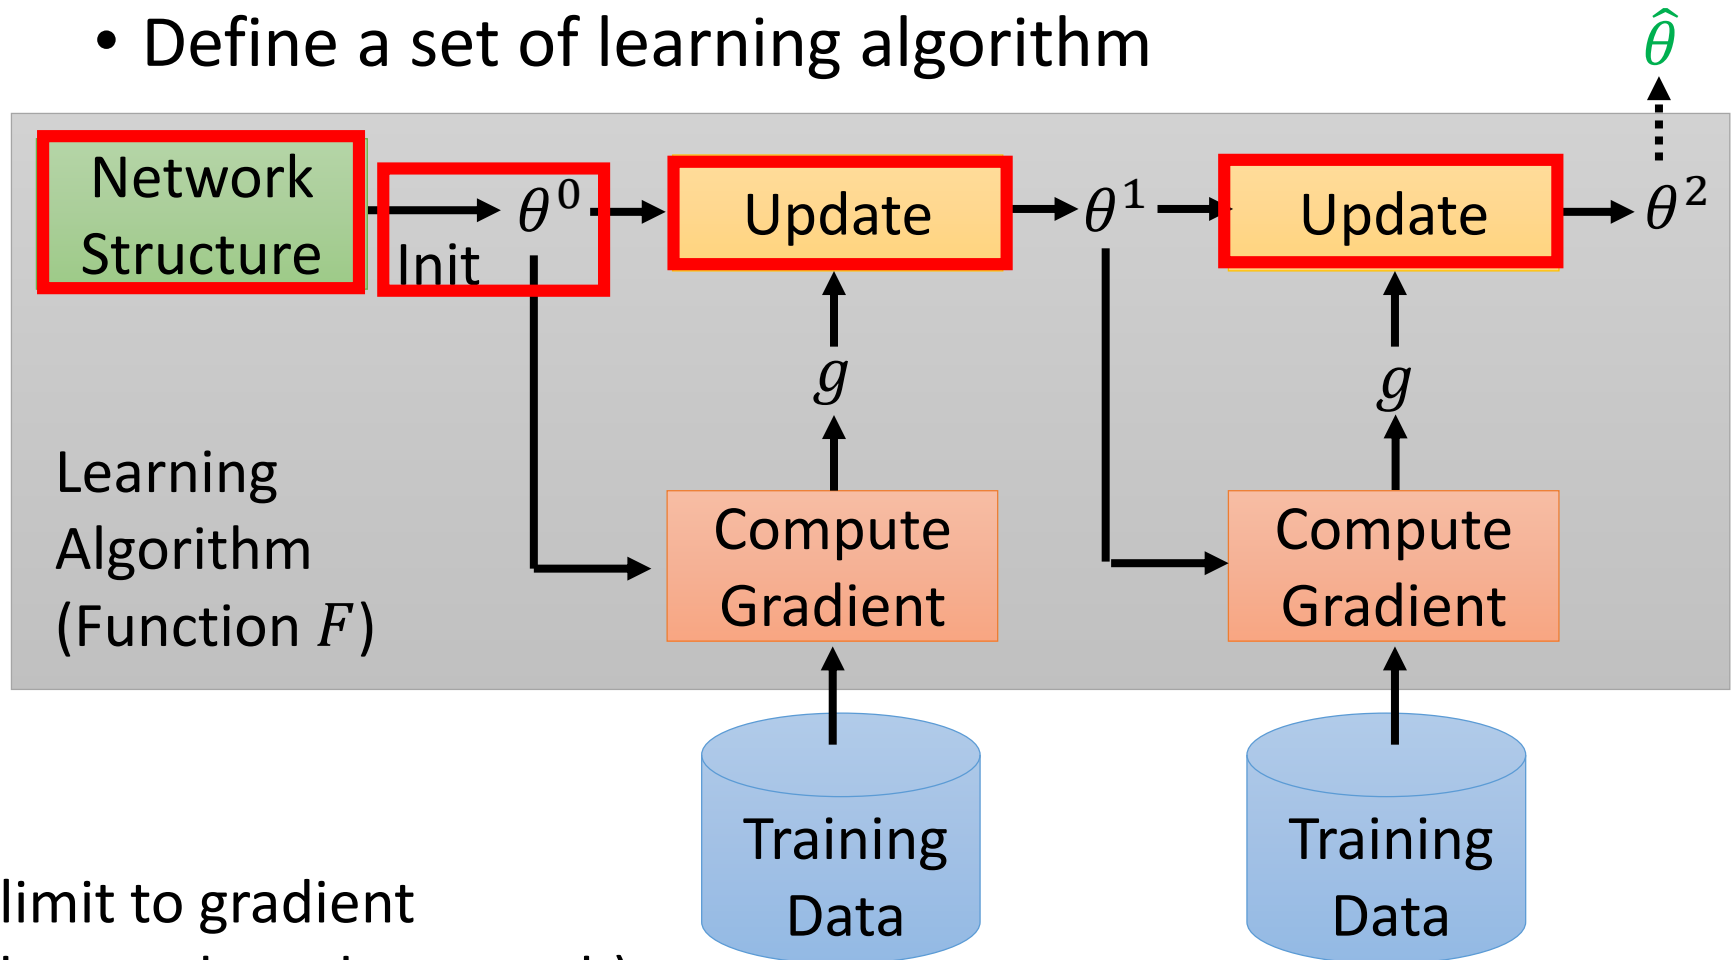
就好像把大象放进冰箱



Meta Learning

Different decisions in the red boxes lead to different algorithms. What happens in the red boxes is decided by humans until now.

- Define a set of learning algorithm



(limit to gradient descent based approach)

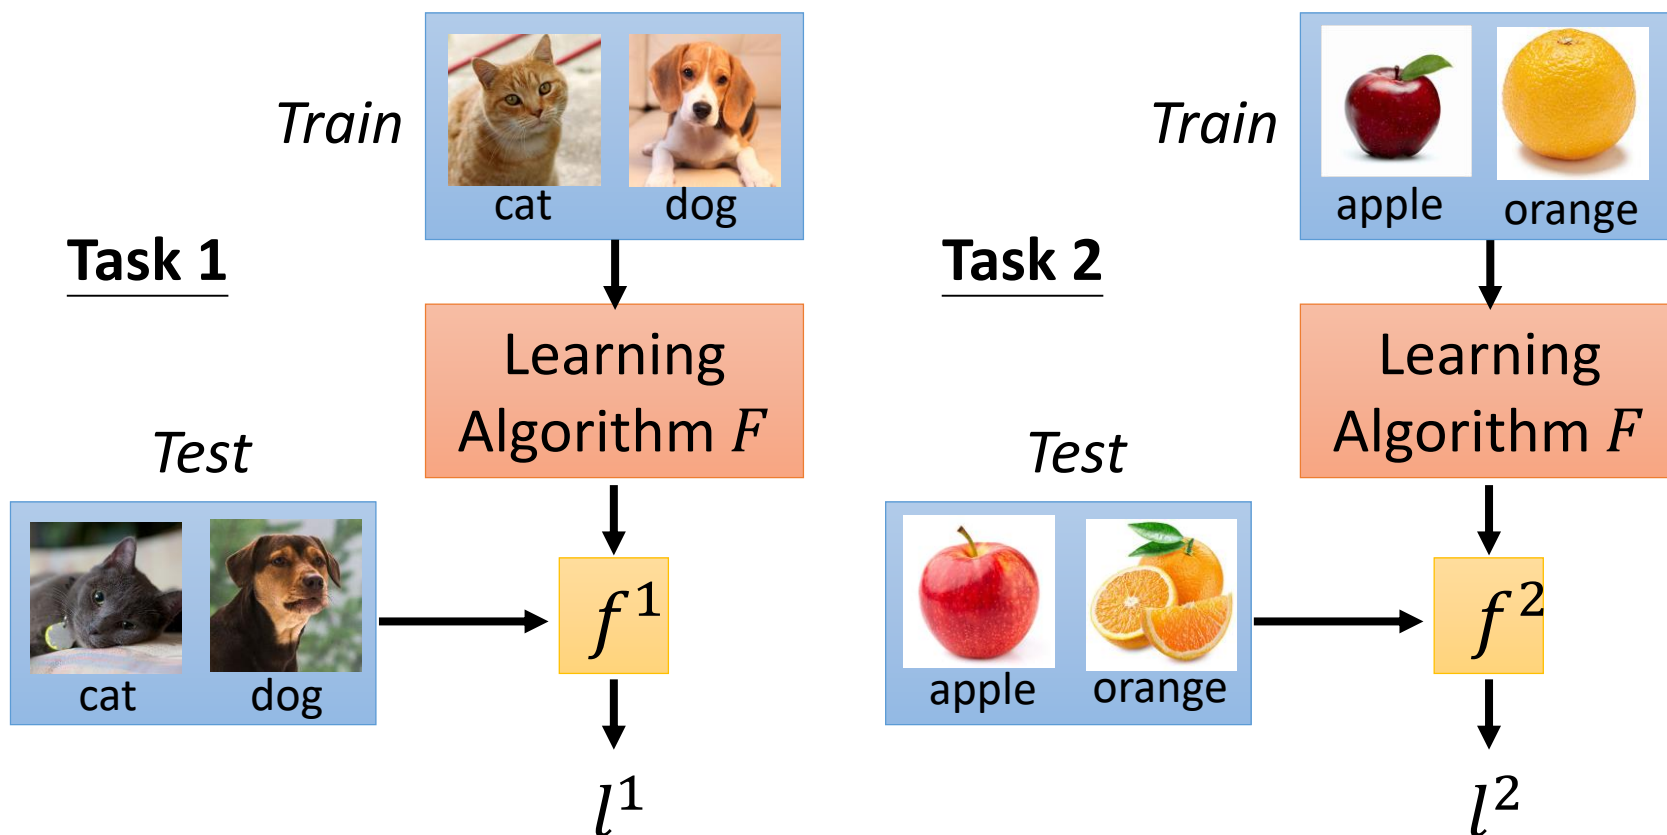
Meta Learning

$$L(F) = \sum_{n=1}^N l^n$$

N → N tasks

l^n → Testing loss for task n after training

- Defining the goodness of a function F



Meta Learning

Widely considered in
few-shot learning

Training
Tasks

Task 1

Support set

Train



cat



dog

Test

Query set



cat



dog

Task 2

Train

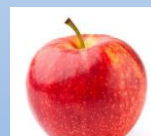


apple



orange

Test



apple

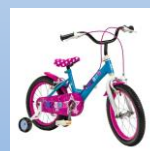


orange

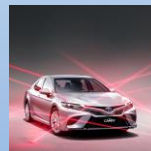
Sometimes you need
validation tasks

Testing Tasks

Train



bike

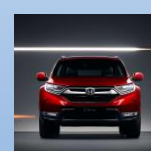


car

Test



bike

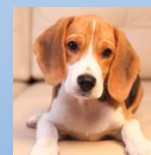


car

Machine Learning



cat



dog

Train



cat



dog

Test

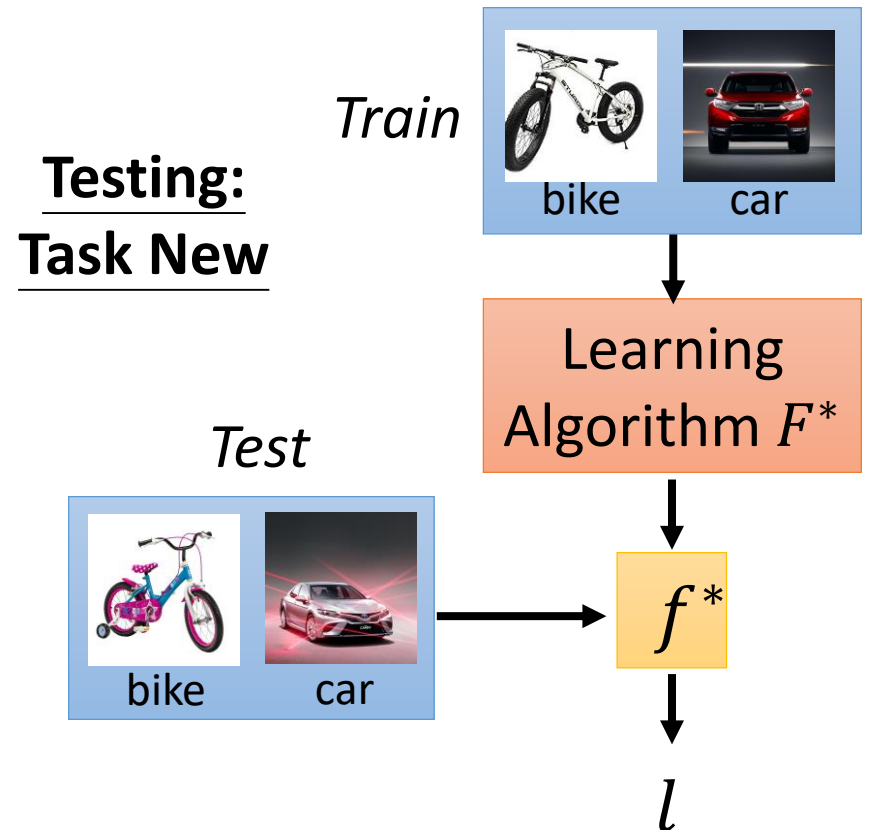
Meta Learning

- Defining the goodness of a function F

$$L(F) = \sum_{n=1}^N l^n$$

- Find the best function F^*

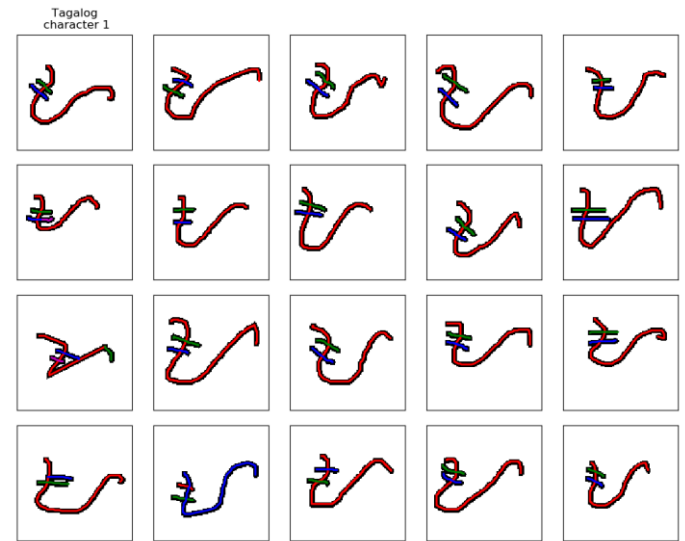
$$F^* = \arg \min_F L(F)$$



Omniglot

<https://github.com/brendenlake/omniglot>

- 1623 characters
- Each has 20 examples



Omniglot

– Few-shot Classification

- **N-ways K-shot** classification: In each training and test tasks, there are **N classes**, each has **K examples**.

20 ways
1 shot

Each character
represents a class

𑖀	𑖁	𑖂	𑖃	𑖄
𑖅	𑖆	𑖇	𑖈	𑖉
𑖊	𑖋	𑖌	𑖍	𑖎
𑖏	𑖐	𑖑	𑖒	𑖓



Testing set
(Query set)

Training set
(Support set)

- Split your characters into training and testing characters
 - Sample N training characters, sample K examples from each sampled characters → one training task
 - Sample N testing characters, sample K examples from each sampled characters → one testing task

Techniques Today

- **MAML**

- Chelsea Finn, Pieter Abbeel, and Sergey Levine, “Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks”, ICML, 2017

- **Reptile**

- Alex Nichol, Joshua Achiam, John Schulman, On First-Order Meta-Learning Algorithms, arXiv, 2018

MAML

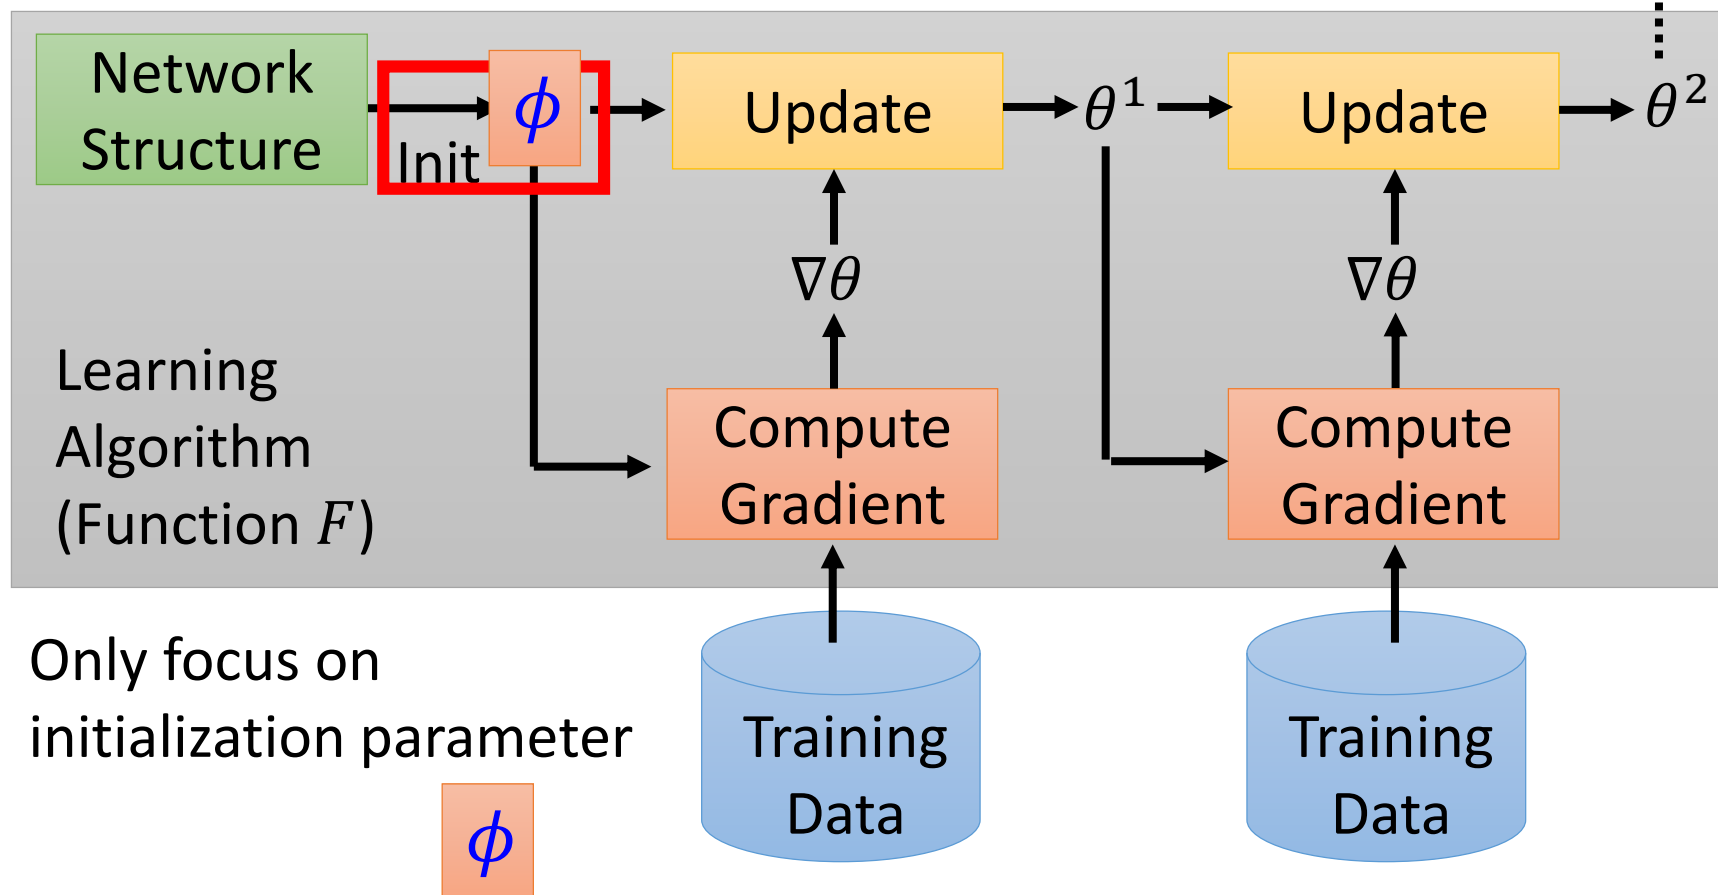
$\hat{\theta}^n$: model learned from task n

Loss Function:

$$L(\phi) = \sum_{n=1}^N l^n(\hat{\theta}^n)$$

$l^n(\hat{\theta}^n)$: loss of task n on the testing set of task n

$\hat{\theta}^n$ depends on ϕ



MAML

$\hat{\theta}^n$: model learned from task n

Loss Function:

$\hat{\theta}^n$ depends on ϕ

$$L(\phi) = \sum_{n=1}^N l^n(\hat{\theta}^n)$$

$l^n(\hat{\theta}^n)$: loss of task n on the testing set of task n

How to minimize $L(\phi)$? Gradient Descent

$$\phi \leftarrow \phi - \eta \nabla_{\phi} L(\phi)$$

Model Pre-training

Widely used in
transfer learning

Loss Function:


$$L(\phi) = \sum_{n=1}^N l^n(\phi)$$

MAML

$$L(\phi) = \sum_{n=1}^N l^n(\hat{\theta}^n)$$

我们不在意 ϕ 在 training task 上表现如何

我们在意用 ϕ 训练出来的 $\hat{\theta}^n$ 表现如何




Model Pre-training

$$L(\phi) = \sum_{n=1}^N l^n(\phi)$$

找寻在所有 task 都最好的 ϕ

并不保证拿 ϕ 去训练以后会得到好的 $\hat{\theta}^n$



MAML

$\hat{\theta}^n$: model learned from task n

Loss Function:

$\hat{\theta}^n$ depends on ϕ

$$L(\phi) = \sum_{n=1}^N l^n(\hat{\theta}^n)$$

$l^n(\hat{\theta}^n)$: loss of task n on the testing set of task n

How to minimize $L(\phi)$? Gradient Descent

$$\phi \leftarrow \phi - \eta \nabla_{\phi} L(\phi)$$

Find ϕ achieving good performance **after training**

潜力

Model Pre-training

Widely used in
transfer learning

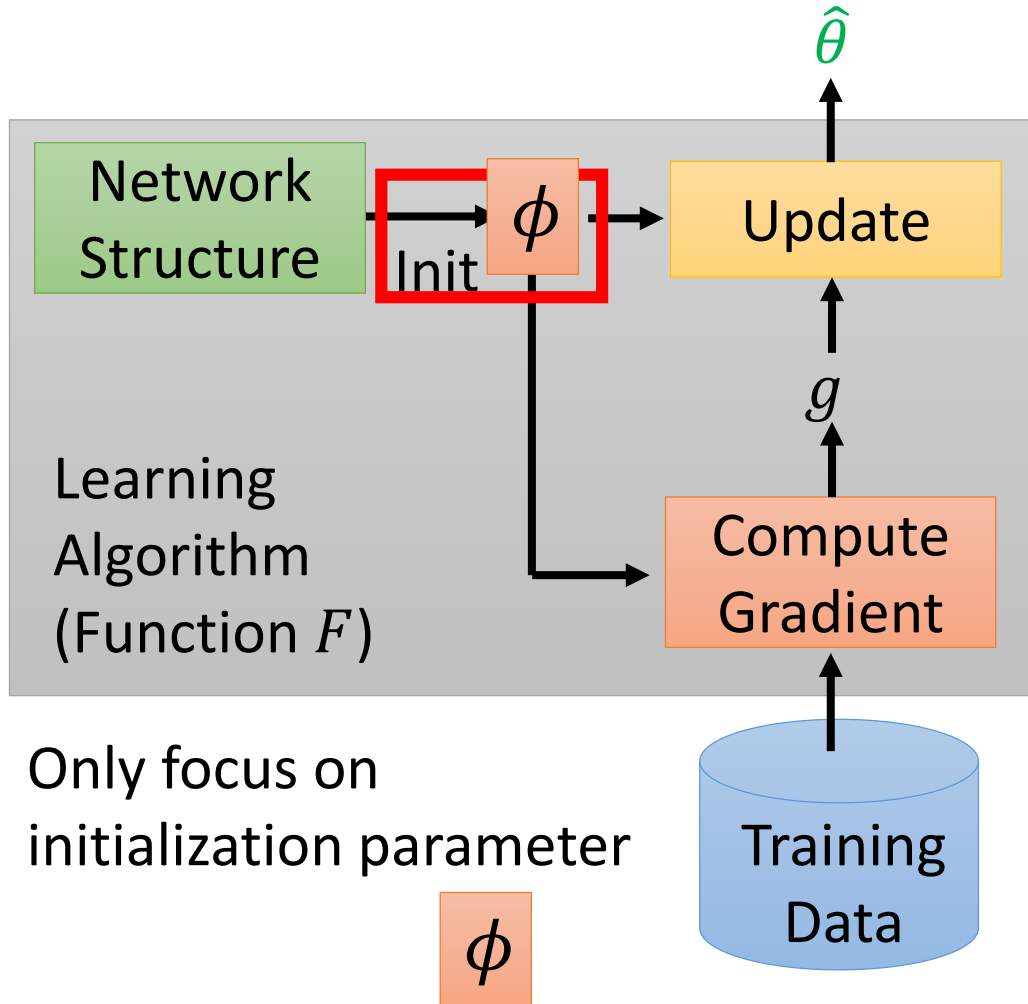
Loss Function:

$$L(\phi) = \sum_{n=1}^N l^n(\phi)$$

Find ϕ achieving good performance

现在表现如何

MAML



$$L(\phi) = \sum_{n=1}^N l^n(\hat{\theta}^n)$$

$$\phi \leftarrow \phi - \eta \nabla_{\phi} L(\phi)$$

Considering one-step training:

$$\hat{\theta} = \phi - \varepsilon \nabla_{\phi} l(\phi)$$

Warning of Math

$$\phi \leftarrow \phi - \eta \nabla_{\phi} L(\phi)$$


$$L(\phi) = \sum_{n=1}^N l^n(\hat{\theta}^n)$$

$$\hat{\theta} = \phi - \varepsilon \nabla_{\phi} l(\phi)$$

$$\nabla_{\phi} L(\phi) = \nabla_{\phi} \sum_{n=1}^N l^n(\hat{\theta}^n) = \sum_{n=1}^N \nabla_{\phi} l^n(\hat{\theta}^n)$$

$$\frac{\partial l(\hat{\theta})}{\partial \phi_i} = \sum_j \frac{\partial l(\hat{\theta})}{\partial \hat{\theta}_j} \frac{\partial \hat{\theta}_j}{\partial \phi_i}$$

$$\nabla_{\phi} l(\hat{\theta}) = \begin{bmatrix} \partial l(\hat{\theta}) / \partial \phi_1 \\ \partial l(\hat{\theta}) / \partial \phi_2 \\ \vdots \\ \partial l(\hat{\theta}) / \partial \phi_i \\ \vdots \end{bmatrix}$$



$$\phi \leftarrow \phi - \eta \nabla_{\phi} L(\phi)$$

$$L(\phi) = \sum_{n=1}^N l^n(\hat{\theta}^n)$$

$$\hat{\theta} = \phi - \varepsilon \nabla_{\phi} l(\phi)$$

$$\nabla_{\phi} L(\phi) = \nabla_{\phi} \sum_{n=1}^N l^n(\hat{\theta}^n) = \sum_{n=1}^N \nabla_{\phi} l^n(\hat{\theta}^n)$$

$$\frac{\partial l(\hat{\theta})}{\partial \phi_i} = \sum_j \frac{\partial l(\hat{\theta})}{\partial \hat{\theta}_j} \frac{\partial \hat{\theta}_j}{\partial \phi_i} \approx \frac{\partial l(\hat{\theta})}{\partial \hat{\theta}_i}$$

$$\hat{\theta}_j = \phi_j - \varepsilon \frac{\partial l(\phi)}{\partial \phi_j}$$

$i \neq j$:

$$\frac{\partial \hat{\theta}_j}{\partial \phi_i} = -\varepsilon \frac{\partial l(\phi)}{\partial \phi_i \partial \phi_j} \approx 0$$

$i = j$:

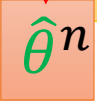
$$\frac{\partial \hat{\theta}_j}{\partial \phi_i} = 1 - \varepsilon \frac{\partial l(\phi)}{\partial \phi_i \partial \phi_j} \approx 1$$

$$\nabla_{\phi} l(\hat{\theta}) = \begin{bmatrix} \partial l(\hat{\theta}) / \partial \phi_1 \\ \partial l(\hat{\theta}) / \partial \phi_2 \\ \vdots \\ \partial l(\hat{\theta}) / \partial \phi_i \\ \vdots \end{bmatrix}$$

$$\phi \leftarrow \phi - \eta \nabla_{\phi} L(\phi)$$

$$L(\phi) = \sum_{n=1}^N l^n(\hat{\theta}^n)$$

$$\hat{\theta} = \phi - \varepsilon \nabla_{\phi} l(\phi)$$

$$\nabla_{\phi} L(\phi) = \nabla_{\phi} \sum_{n=1}^N l^n(\hat{\theta}^n) = \sum_{n=1}^N \nabla_{\phi} l^n(\hat{\theta}^n)$$



$$\frac{\partial l(\hat{\theta})}{\partial \phi_i} = \sum_j \frac{\partial l(\hat{\theta})}{\partial \hat{\theta}_j} \frac{\partial \hat{\theta}_j}{\partial \phi_i} \approx \frac{\partial l(\hat{\theta})}{\partial \hat{\theta}_i}$$

products, which is supported by standard deep learning libraries such as TensorFlow (Abadi et al., 2016). In our experiments, we also include a comparison to dropping this backward pass and using a first-order approximation, which we discuss in Section 5.2.


$$\nabla_{\phi} l(\hat{\theta}) = \begin{bmatrix} \partial l(\hat{\theta}) / \partial \phi_1 \\ \partial l(\hat{\theta}) / \partial \phi_2 \\ \vdots \\ \partial l(\hat{\theta}) / \partial \phi_i \\ \vdots \end{bmatrix} = \begin{bmatrix} \partial l(\hat{\theta}) / \partial \hat{\theta}_1 \\ \partial l(\hat{\theta}) / \partial \hat{\theta}_2 \\ \vdots \\ \partial l(\hat{\theta}) / \partial \hat{\theta}_i \\ \vdots \end{bmatrix} = \nabla_{\hat{\theta}} l(\hat{\theta})$$

End of Warning

MAML – Real Implementation




Model Pre-training




Translation

18 training tasks: 18 different languages translating to English
2 validation tasks: 2 different languages translating to English



Ro = Romanian



Fi = Finnish

<https://arxiv.org/abs/1808.08437>

Techniques Today

- **MAML**


- Chelsea Finn, Pieter Abbeel, and Sergey Levine, “Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks”, ICML, 2017

- **Reptile**

- Alex Nichol, Joshua Achiam, John Schulman, On First-Order Meta-Learning Algorithms, arXiv, 2018

<https://openai.com/blog/reptile/>

Reptile



You might be thinking “isn’t this the same as training on the expected loss $\mathbb{E}_{\tau} [L_{\tau}]$?” and then checking if the date is April 1st. Indeed, if the partial minimization consists of a single gradient step, then this algorithm corresponds to minimizing the expected loss:

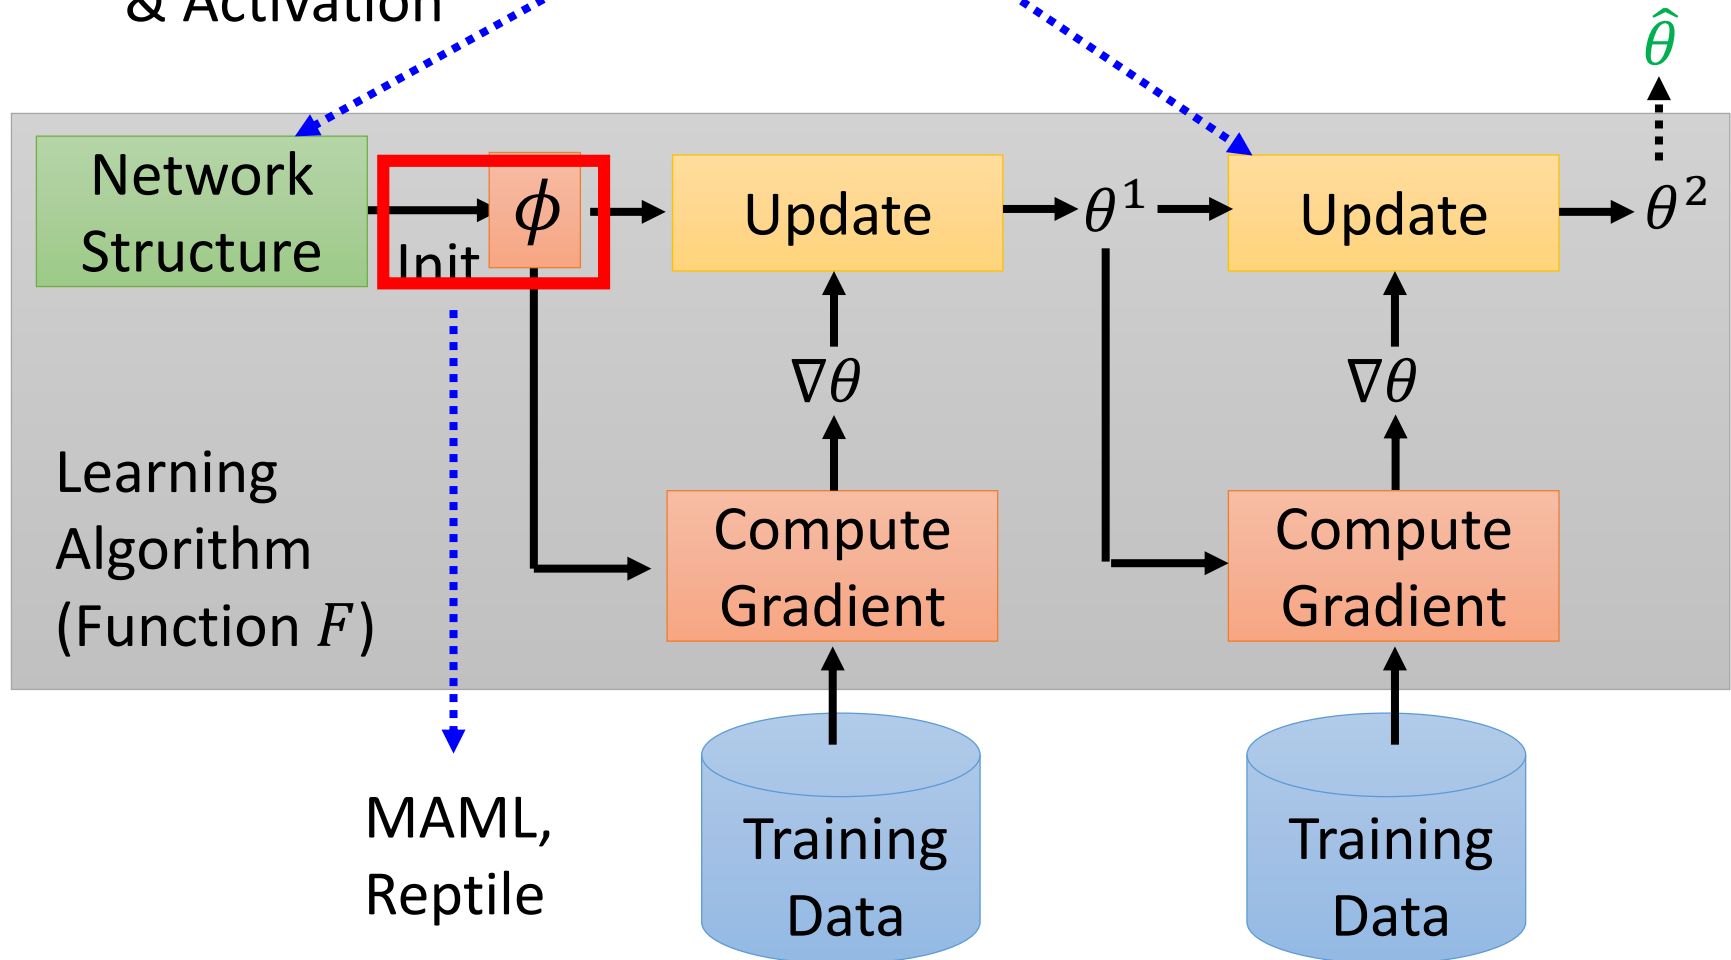
(this sentence is removed in the updated version)

More ... Video: <https://www.youtube.com/watch?v=c10nxBcSH14>

Training a network (by RL) to determine ...

Architecture
& Activation

How to update



Turtles all the way down ?



- We learn the initialization parameter ϕ by gradient descent
 - What is the initialization parameter ϕ^0 for initialization parameter ϕ ?
- Learn
- Learn to learn
- Learn to learn to learn

Crazy Idea?

- How about learning algorithm beyond gradient descent?

